

# End to End: Part 4

Earl Wong

# Subjects

- Optics
- Sensor
- ISP
- **GPU**
- NPU

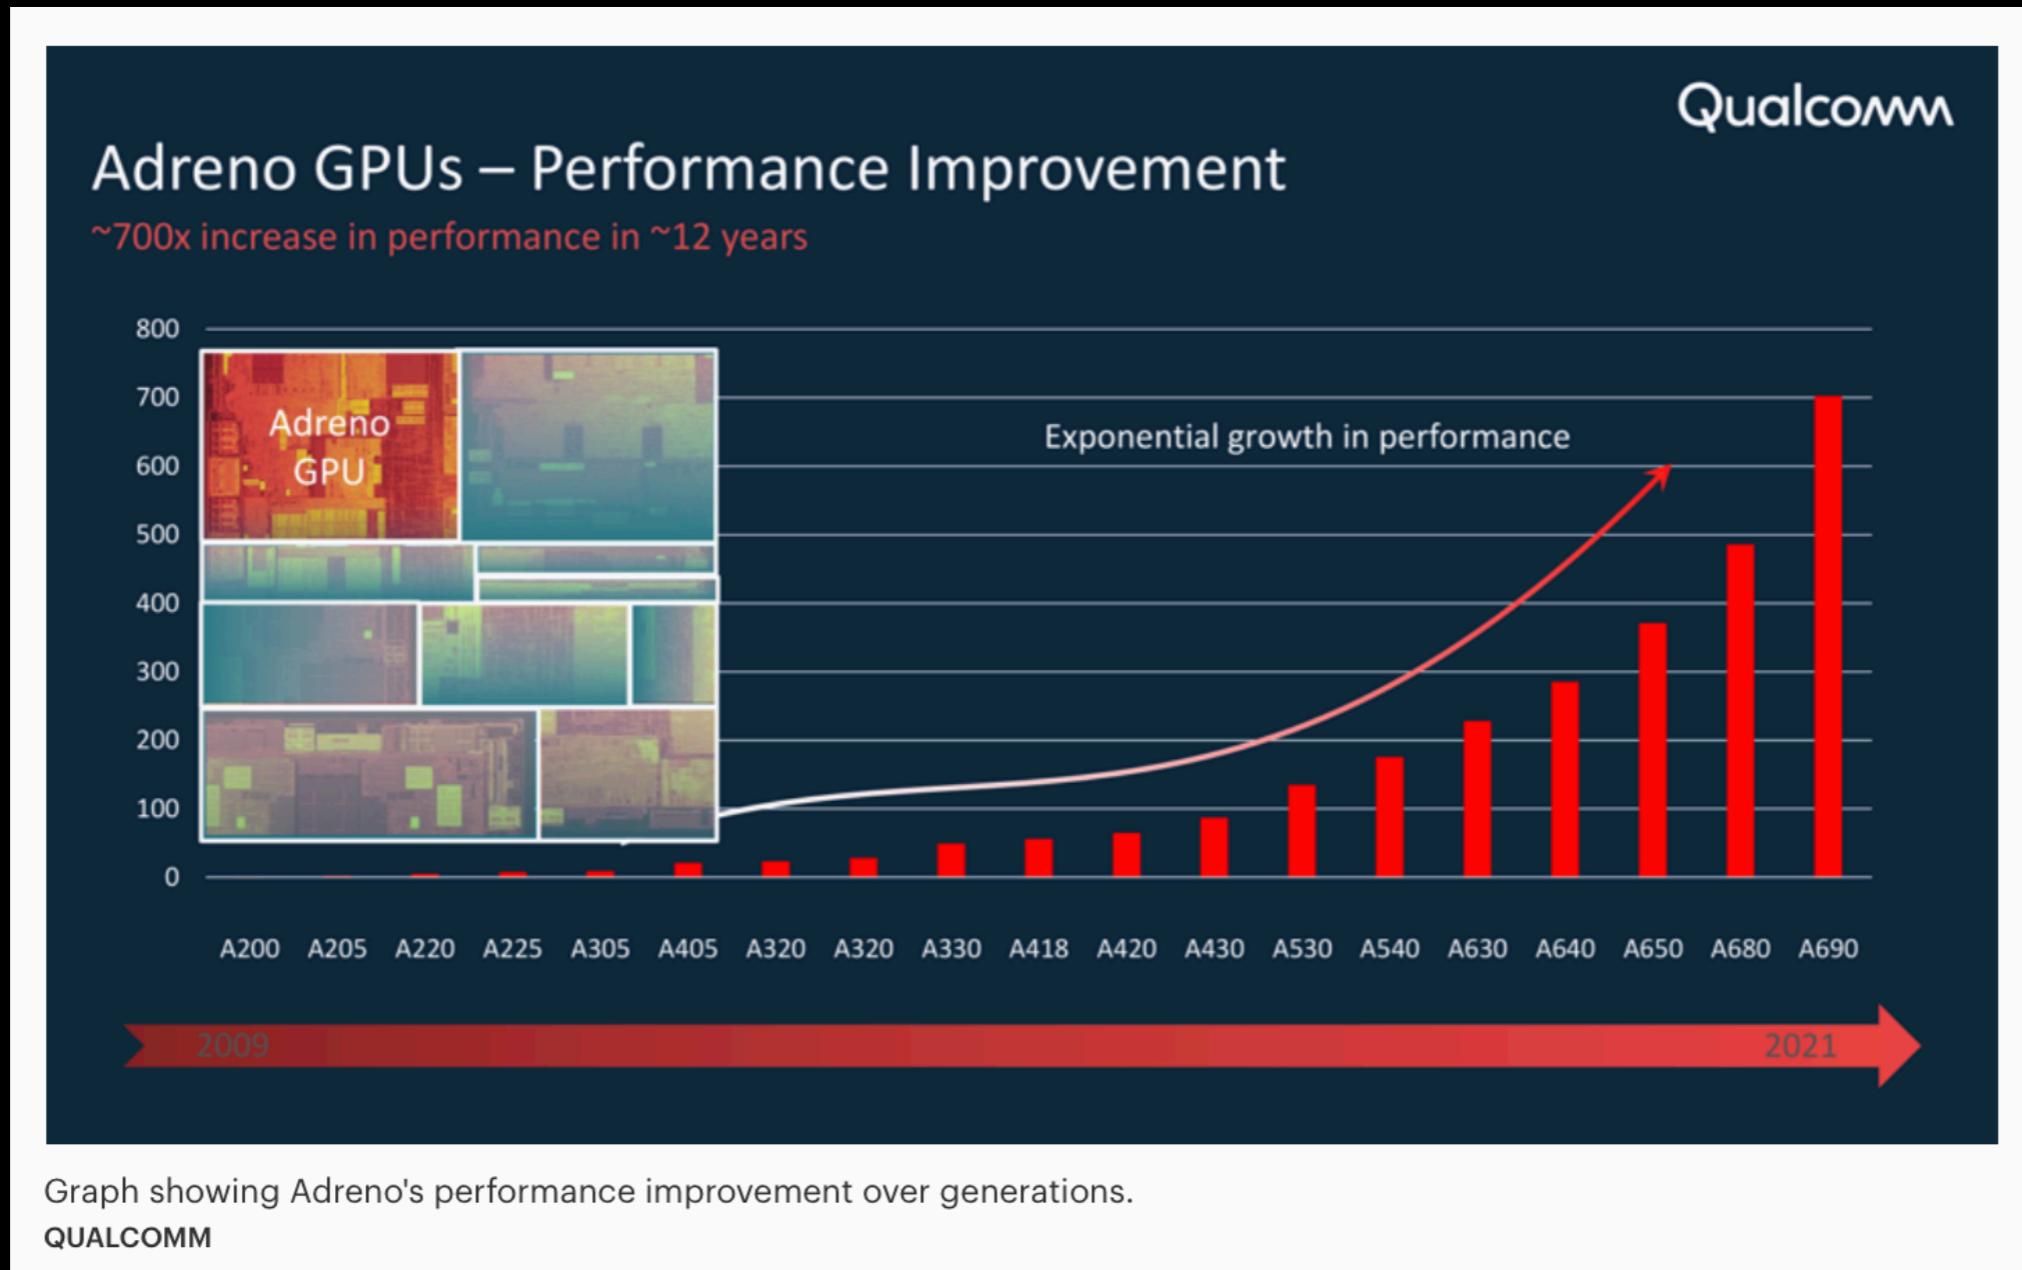
# GPU

- The GPU on mobile devices can handle both image / video post processing and ML processing (**Transformers** and CNN based models).
- The NPU on mobile devices is highly optimized for ML processing - specifically, CNN based models.
- We will discuss the NPU in the next presentation.

# GPU

- We will focus on the Qualcomm (QC) Adreno GPU, since Qualcomm is a major chip supplier for mobile devices.
- The Mediatek Mali GPU provides similar functionality for Mediatek mobile chipsets.
- Note: Micro level architectural descriptions for the QC Adreno GPU hardware do not exist.
- Hence, the Compute Unit (CU) and Shader Processor (SP) architectural diagrams are my “best guesses”, based on Snapdragon Profiler output fields, known approaches from other GPU manufacturers, etc.

# GPU Evolution



2024 to 2025 Flagship QC GPU - Adreno 800 Series

# GPU - OpenGL support

**AMD**  
Mobile Acquisition  
2009

Pre- 2006  
Software Graphics  
Rendering  
(with CPU/DSP)

2014-2015  
Adreno 4x Series GPUs  
(DX11.2, HW TS, GS, OpenCL Full Profile, etc.)

2012-14  
Adreno 3x Series GPUs  
(First OpenGL ES 3.0 capable GPUs, GPGPU accelerated  
OpenCL Support, 100% in-house developed)

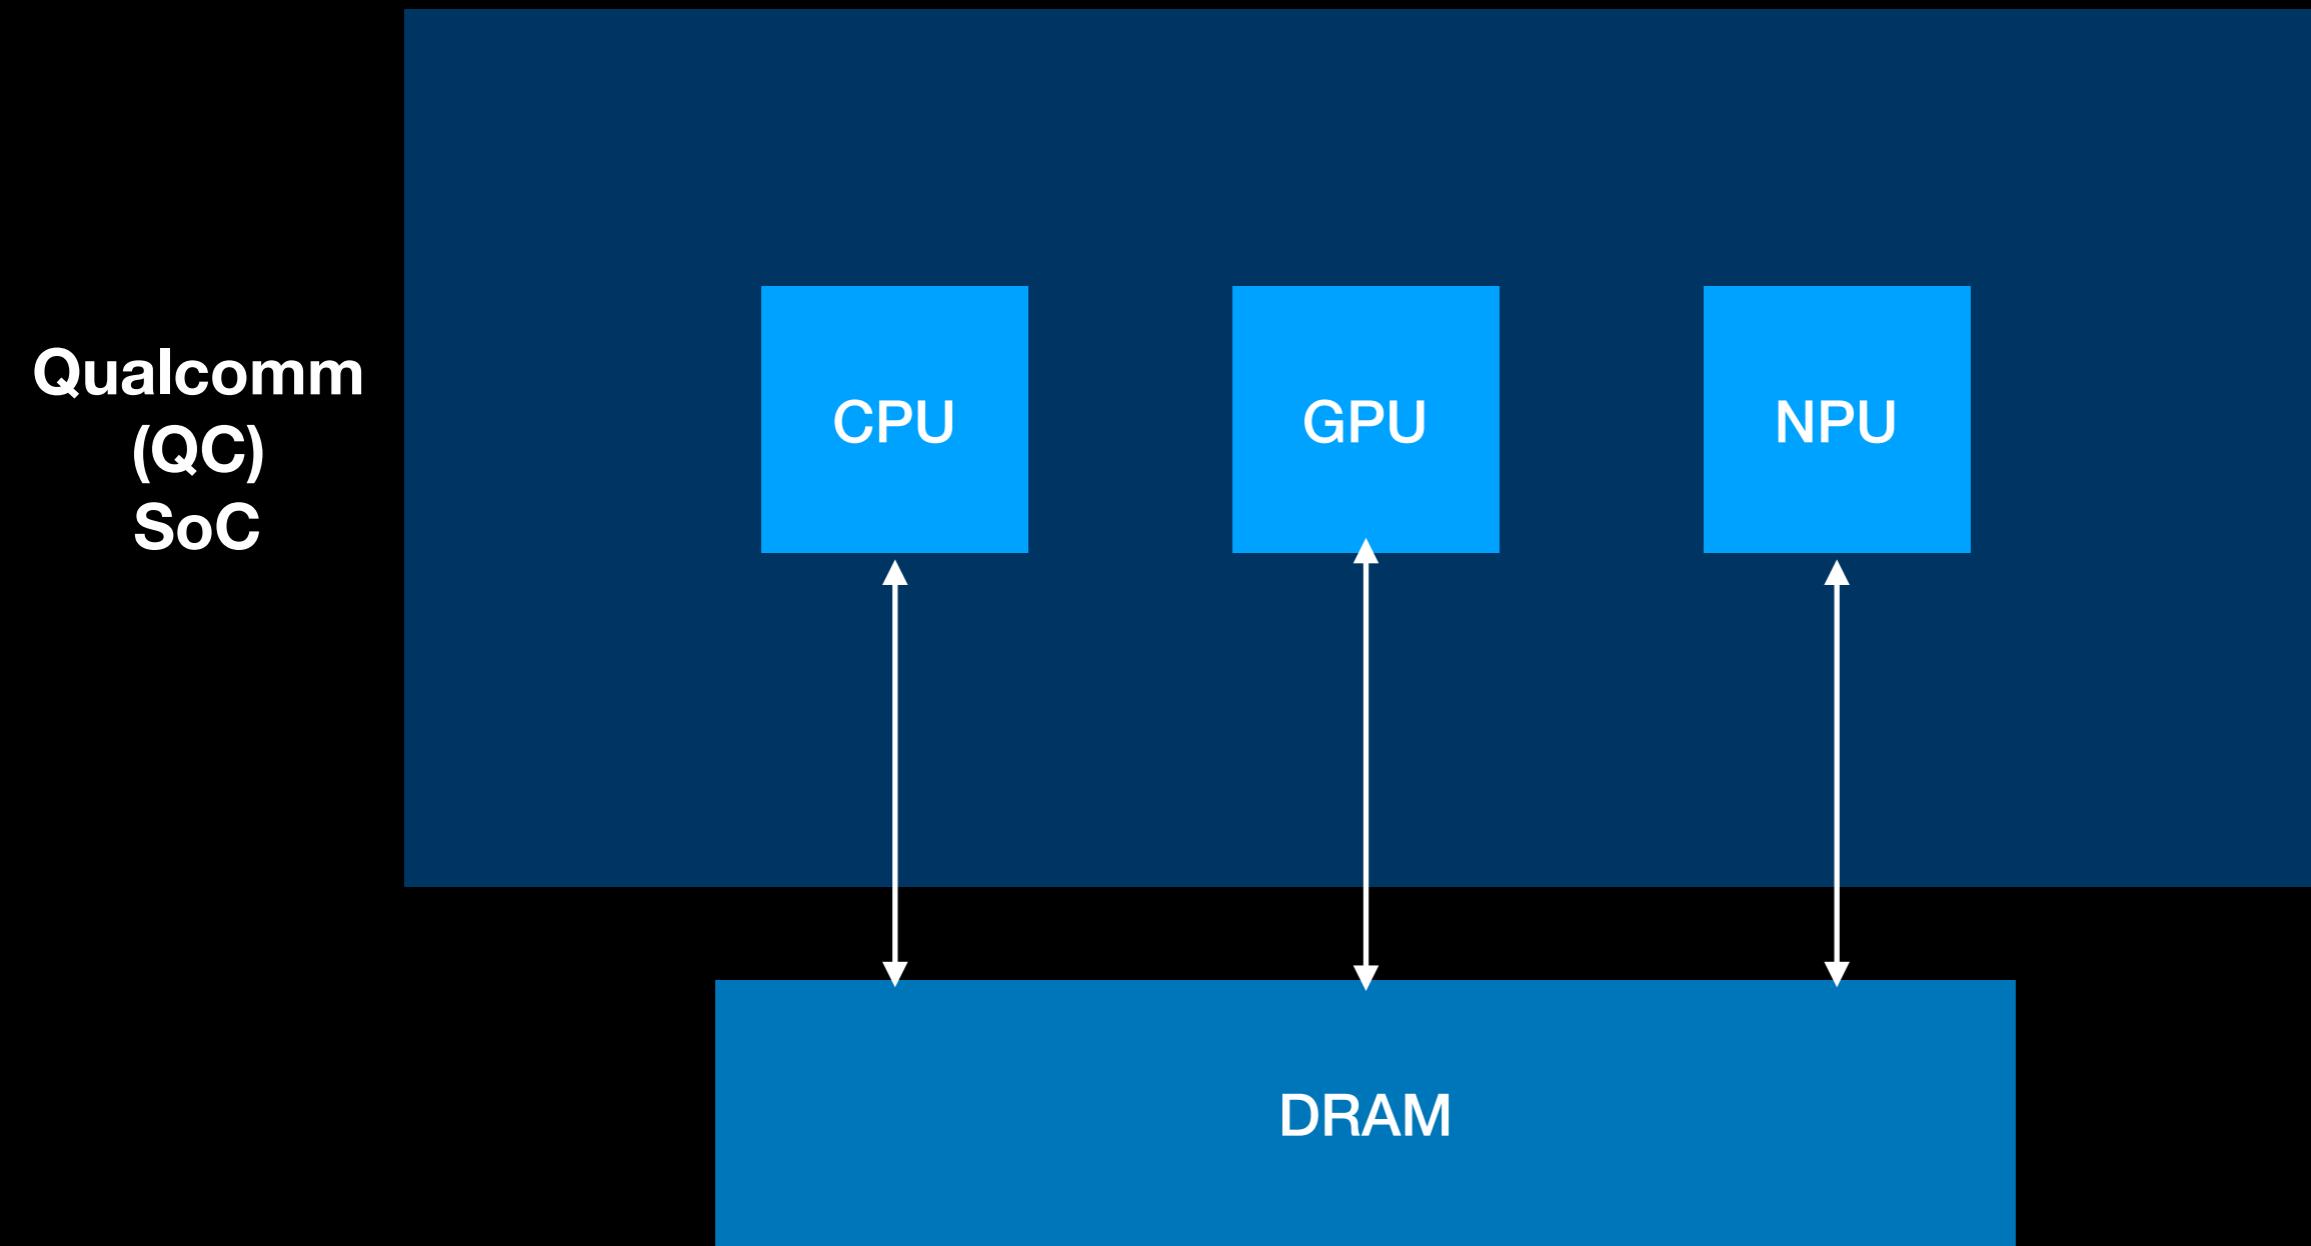
2009-2011  
Adreno 2x Series GPUs  
(Fully programmable OpenGL ES 2.0 GPUs,  
A200 licensed, A205 jointly developed, A220/225 in house)

2006 - 2008  
Adreno 1x Series GPUs  
(Fixed Function GPU hw acceleration, A110 & A120 in  
house developed, A130 licensed from ATI/AMD)



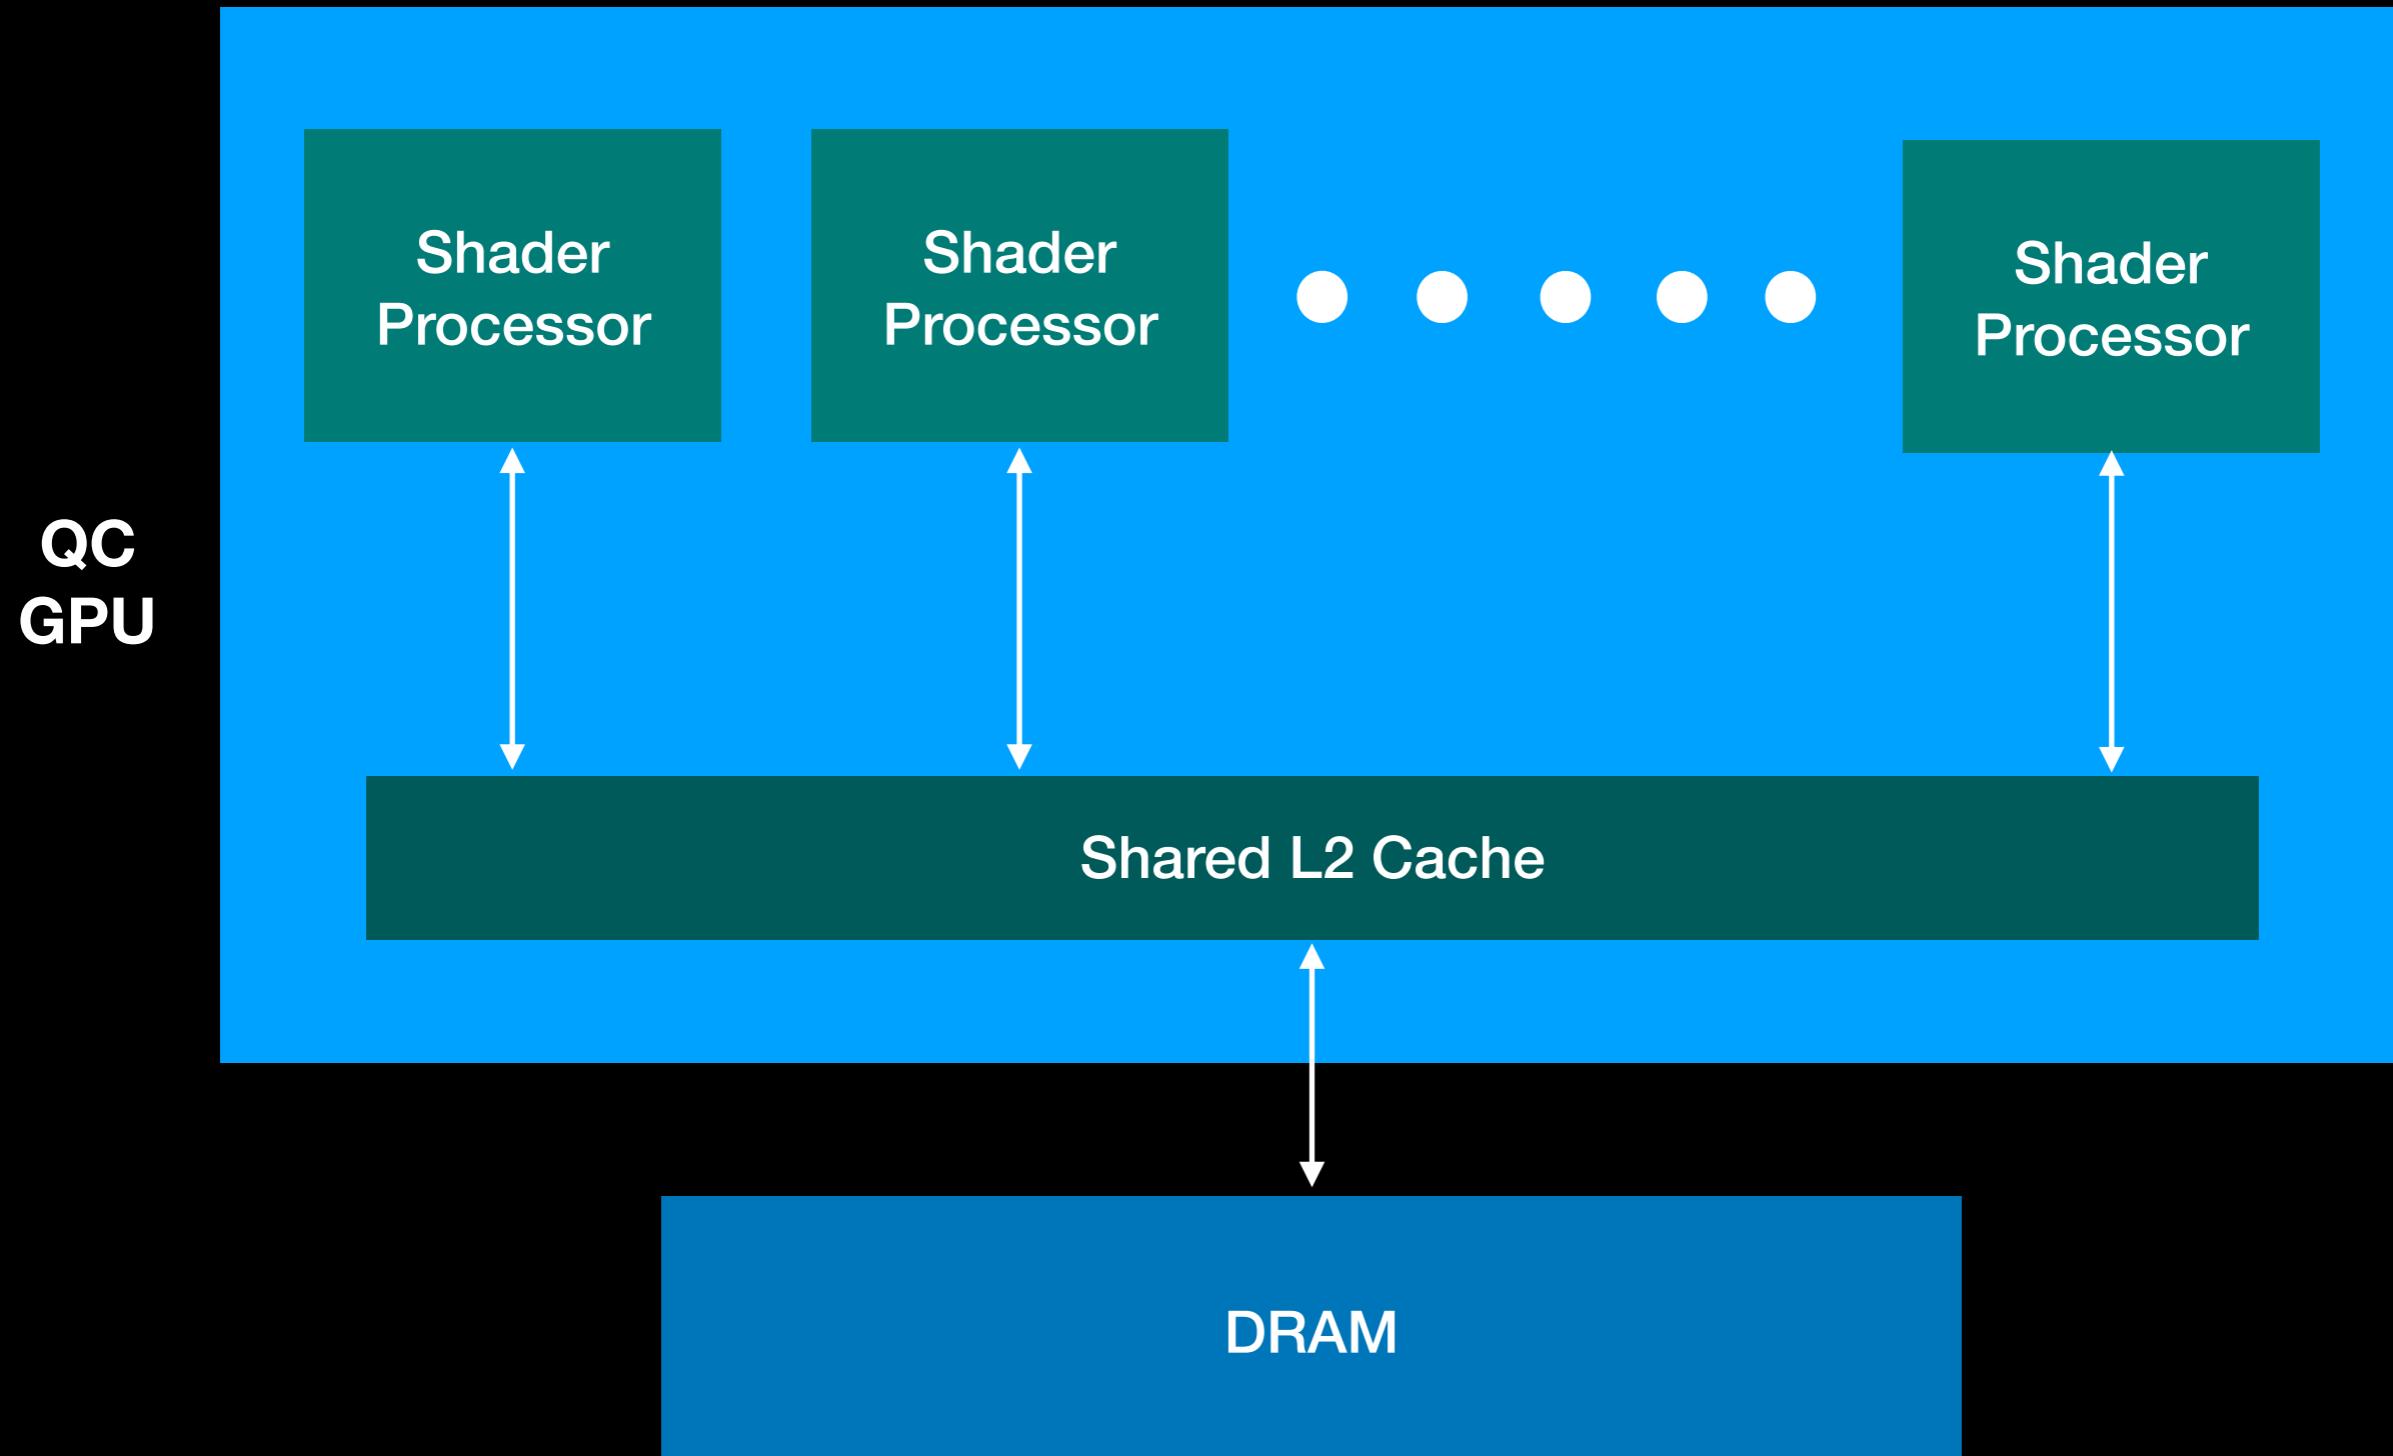
OpenGL ES 3.2: Full support in the 600 series to current (800 series)

# SoC View

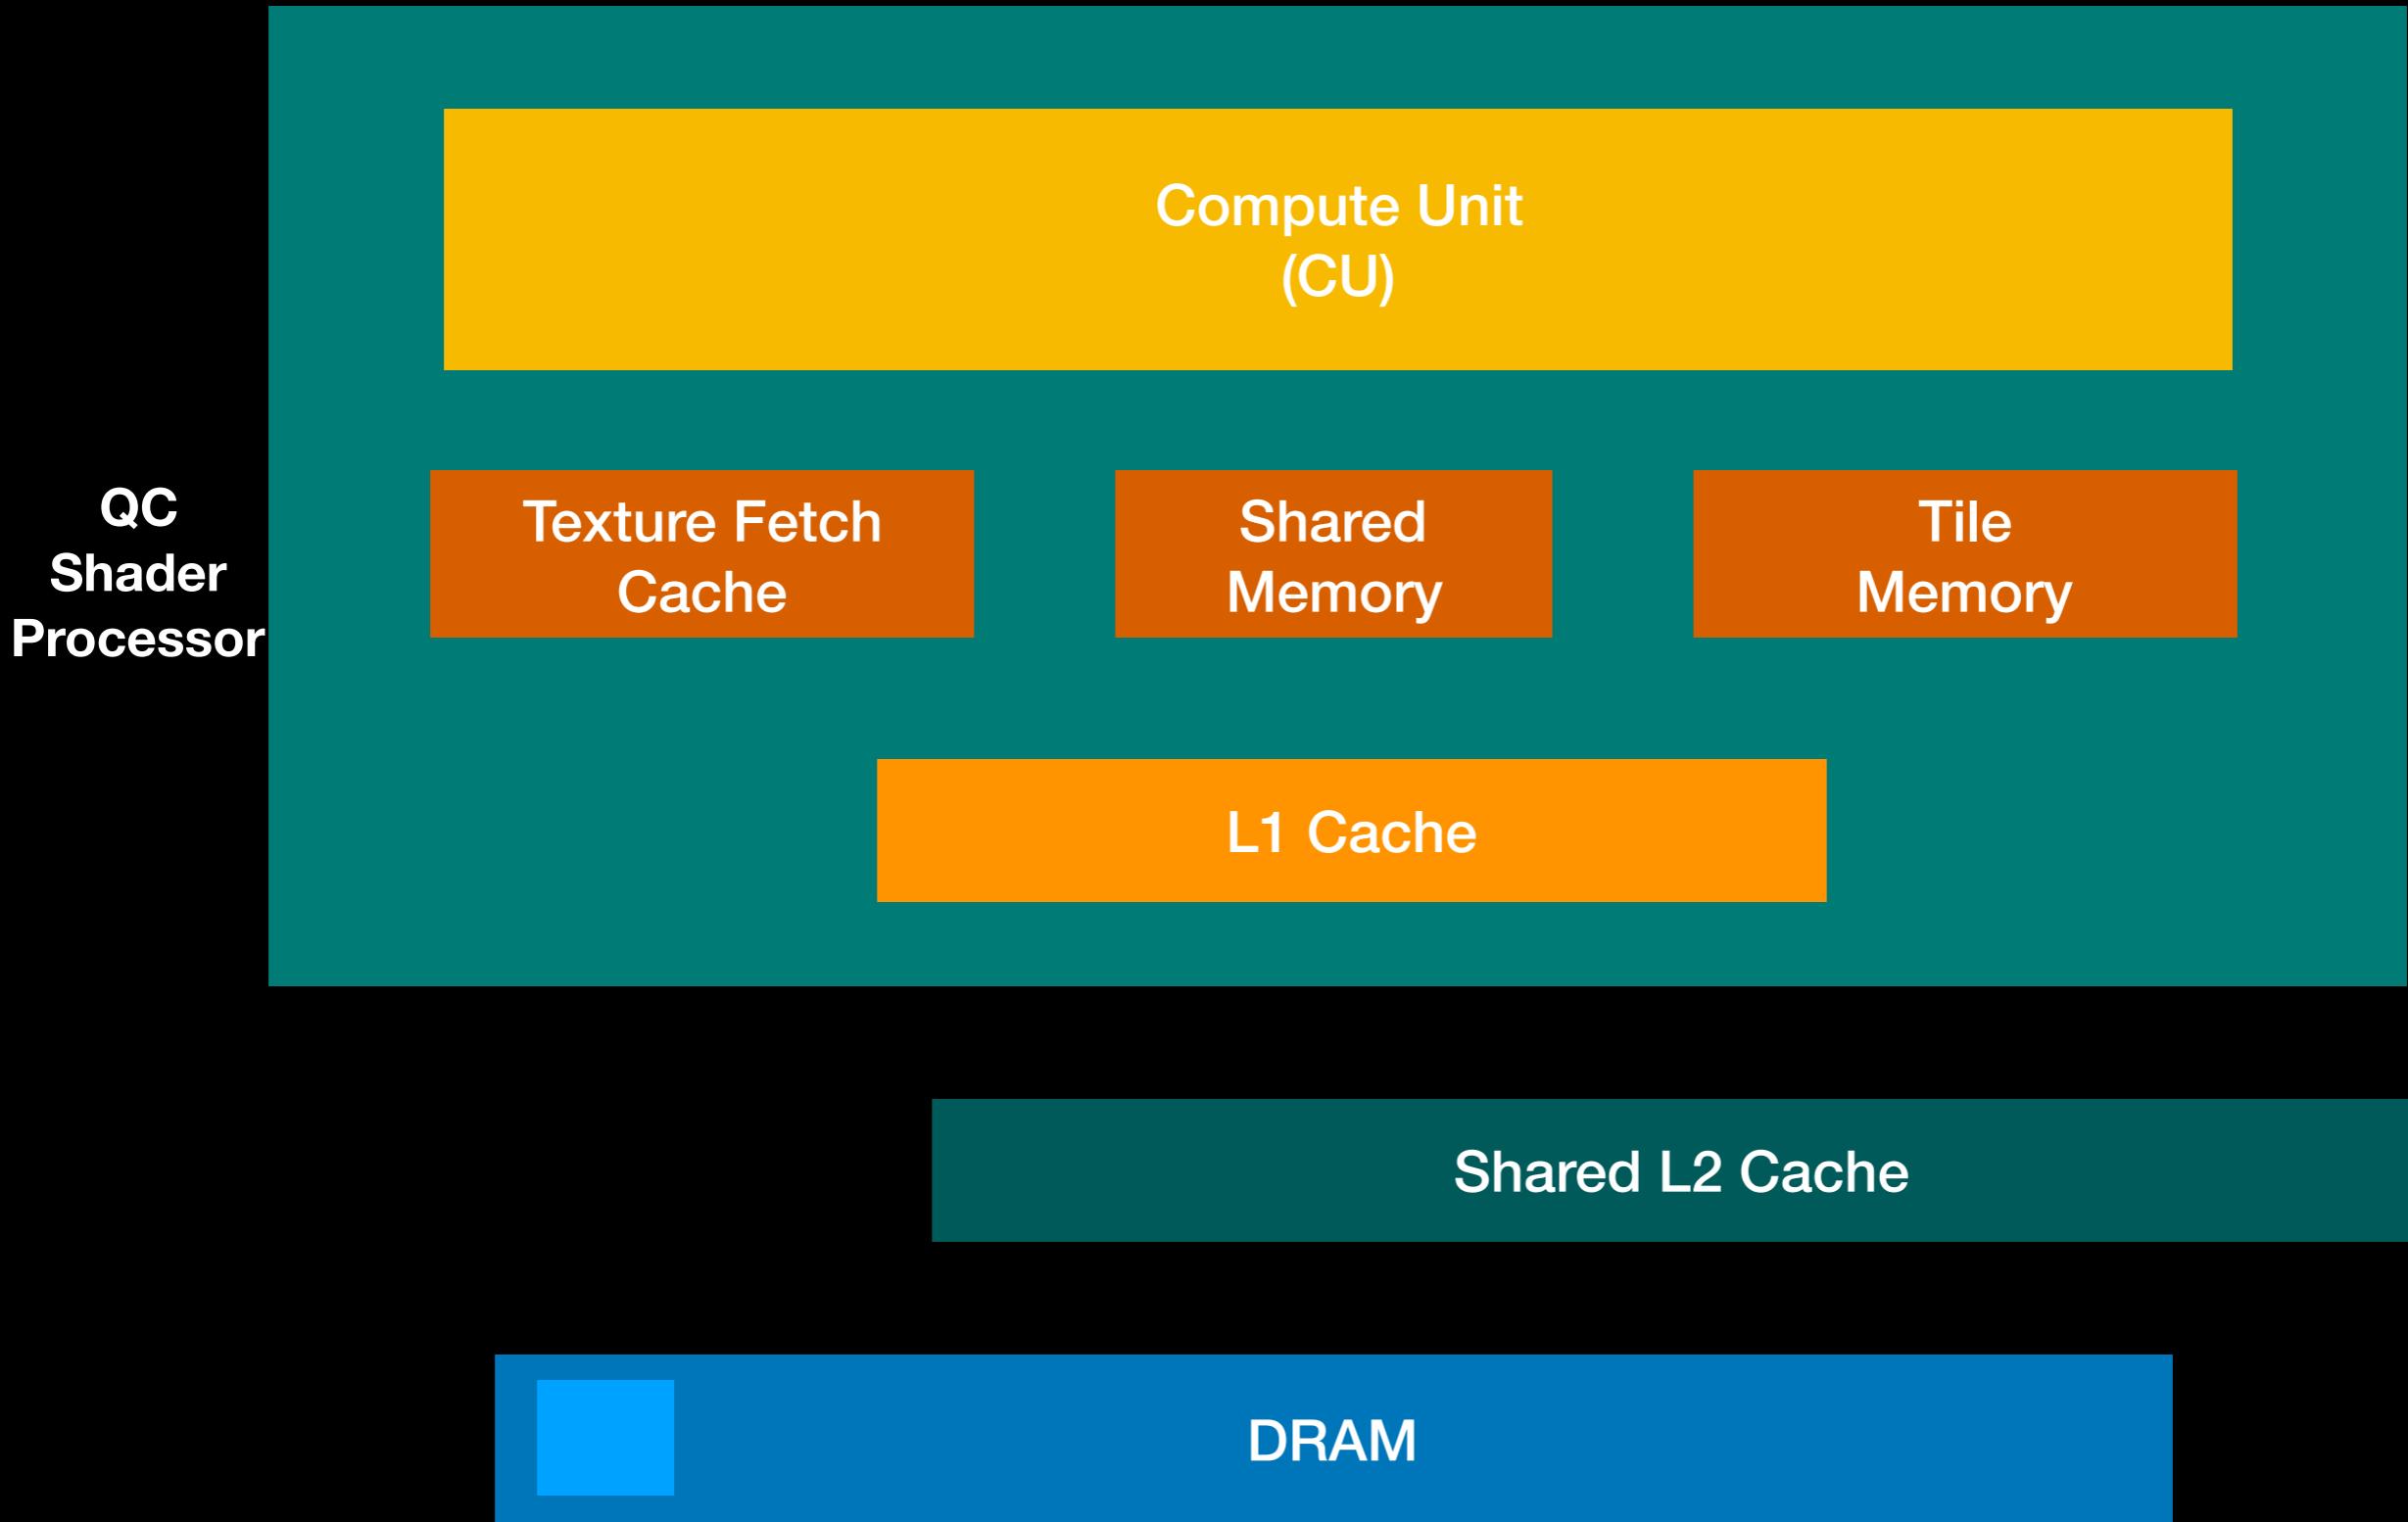


**DRAM is LPDDR# - low power double data rate RAM for mobile devices.  
[Clocking occurs on both the rising and falling edges.]**

# GPU View



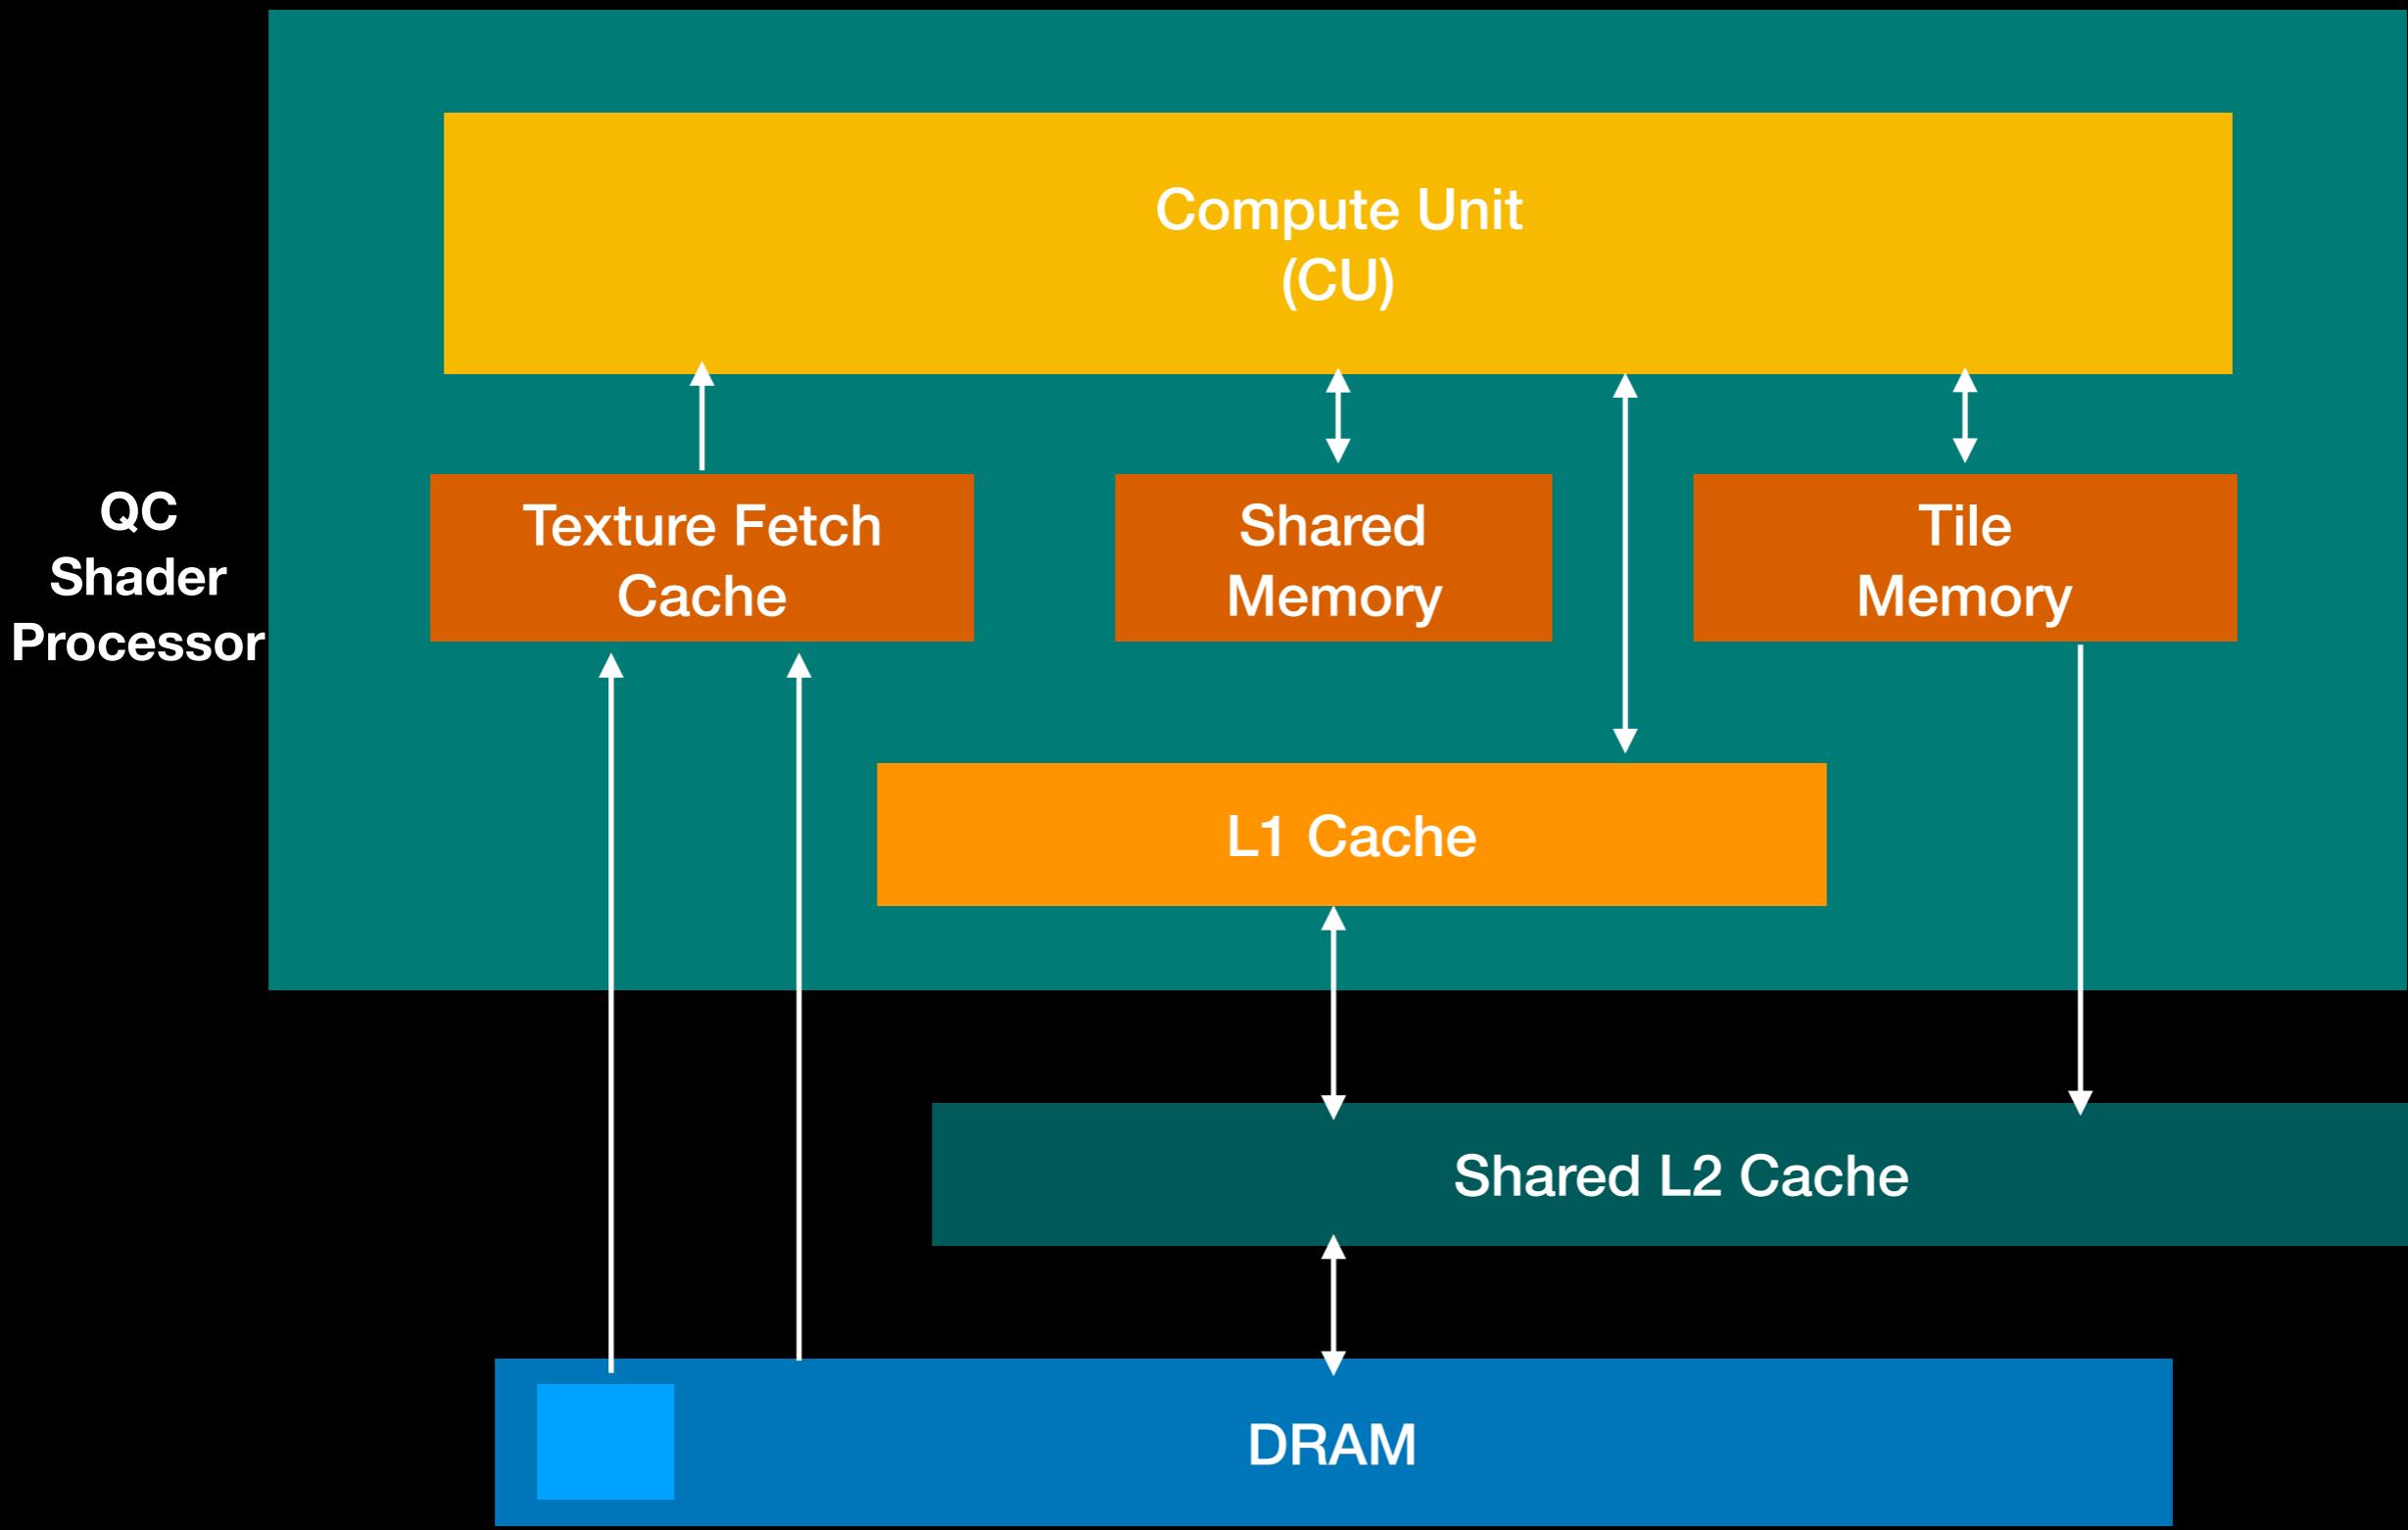
# Shader Processor View



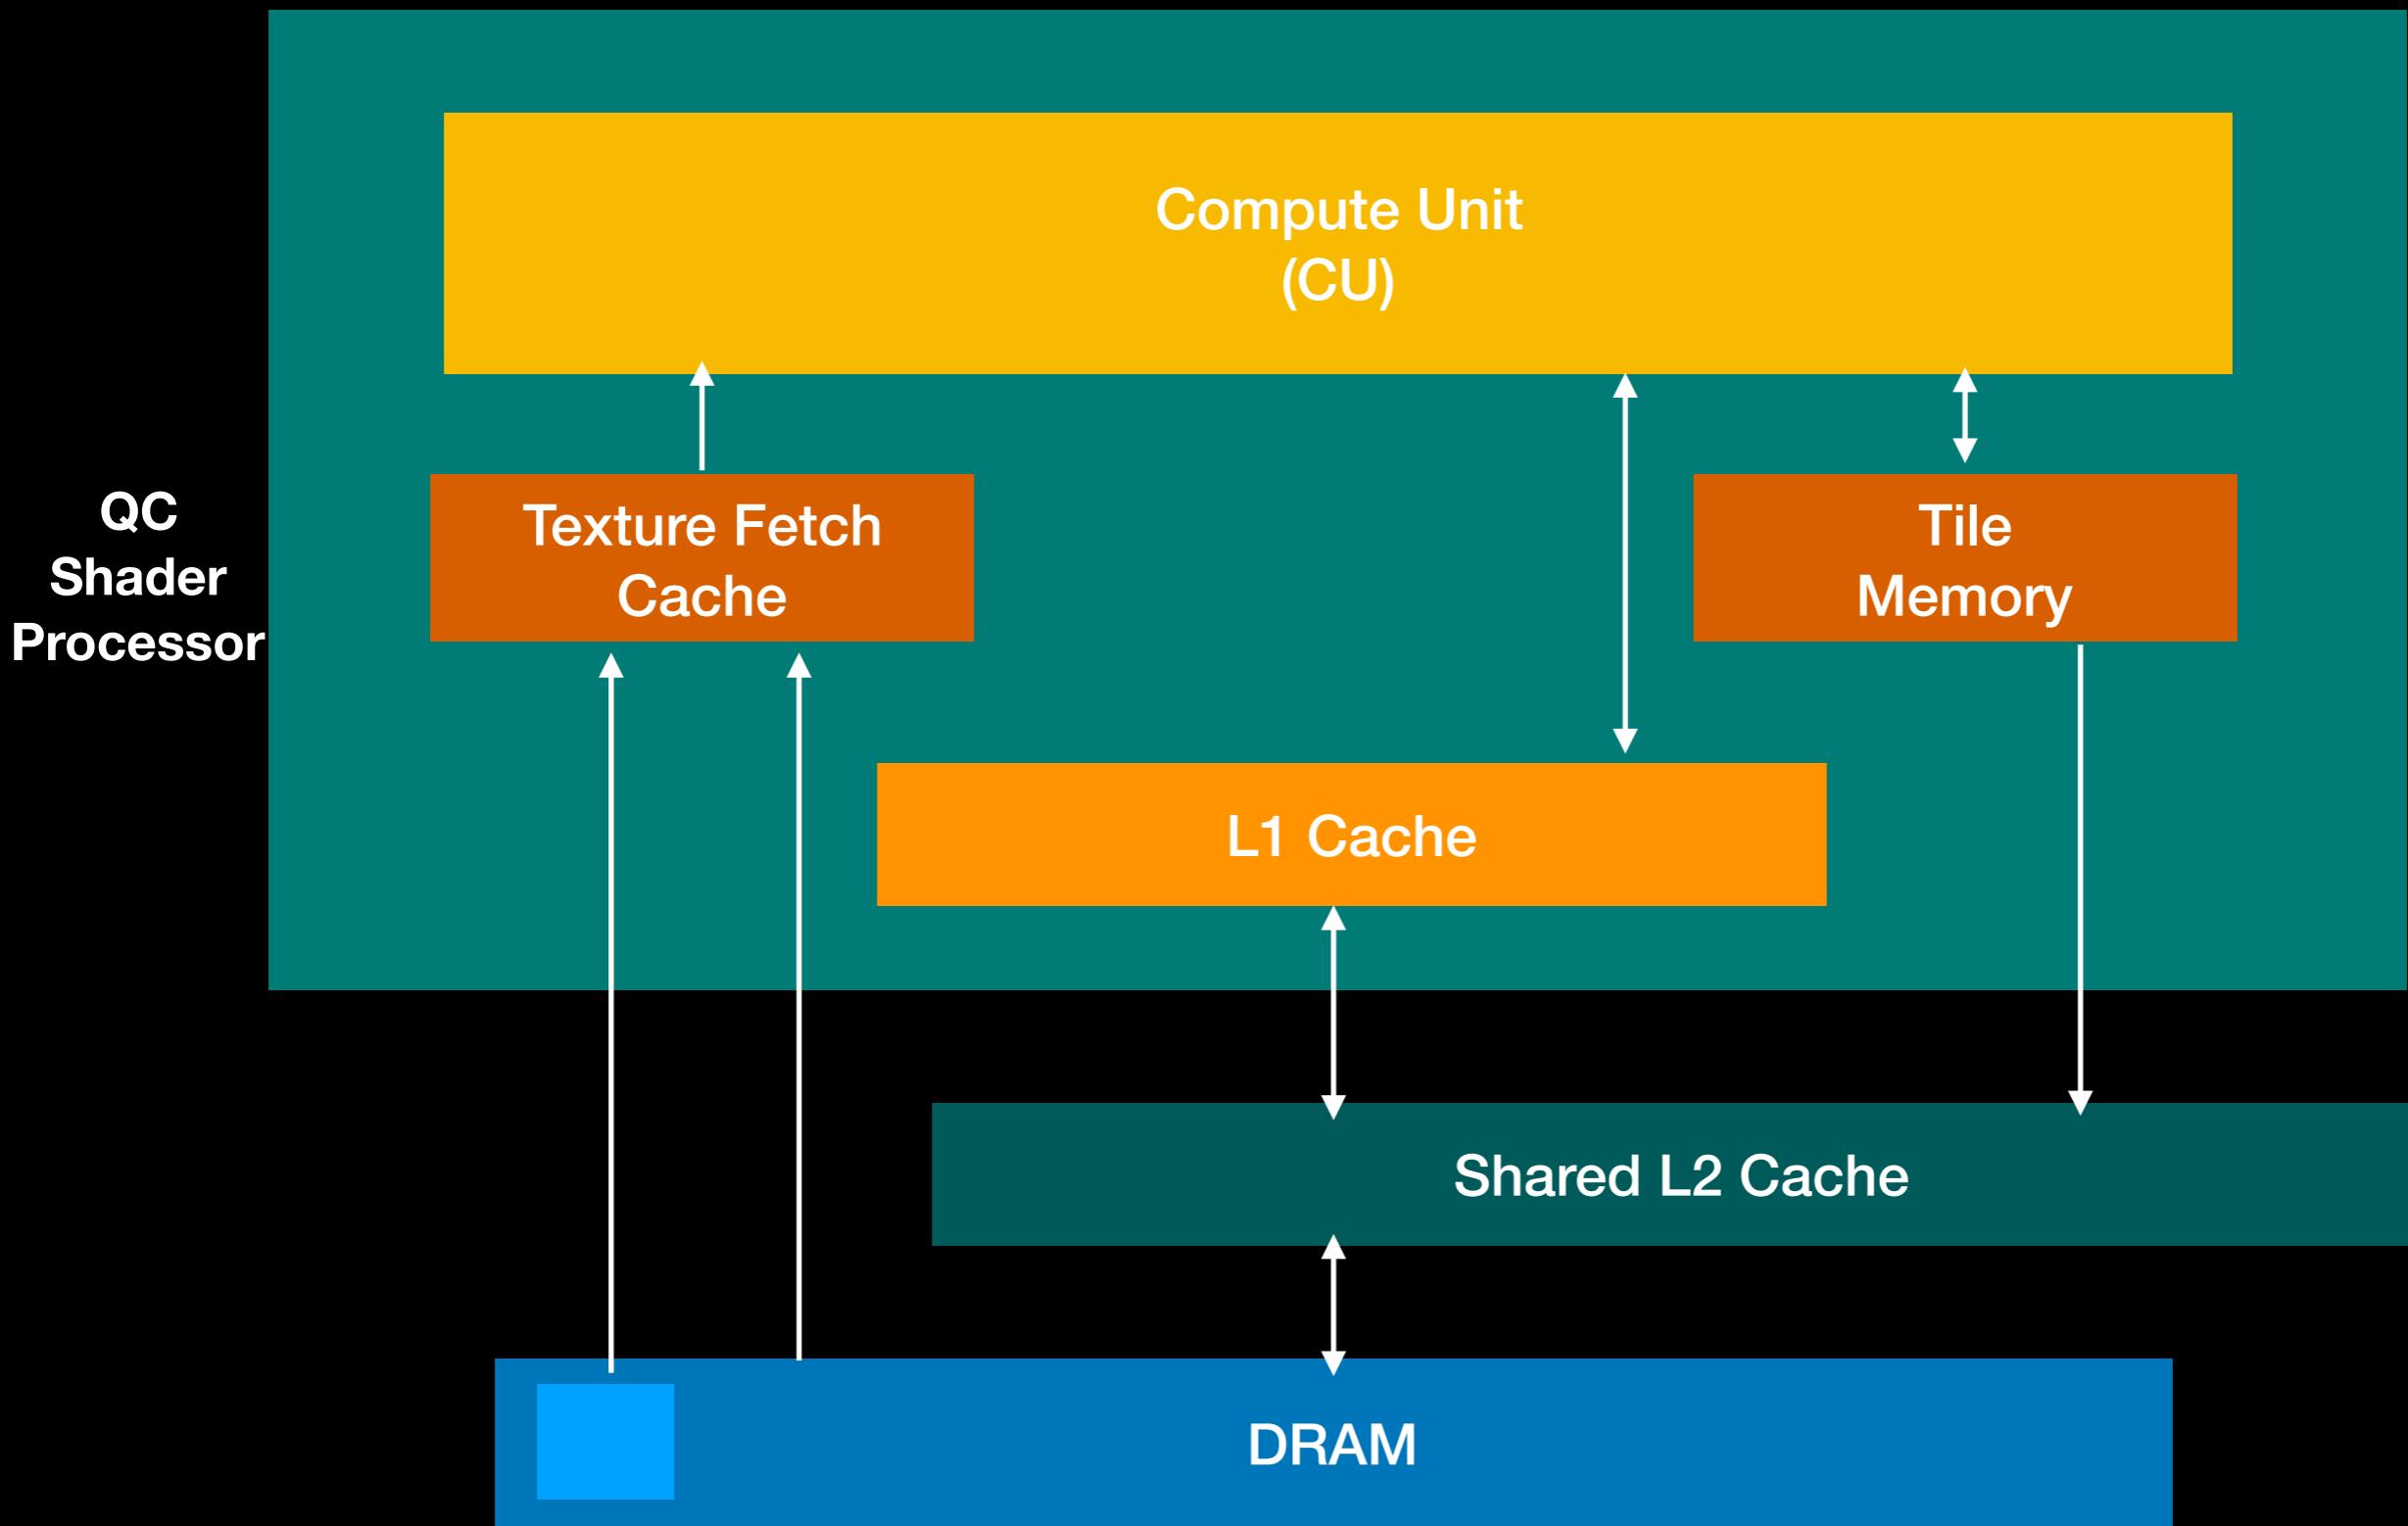
# Shader Process View

- I will now add “the arrows”.
- Caveat emptor ....

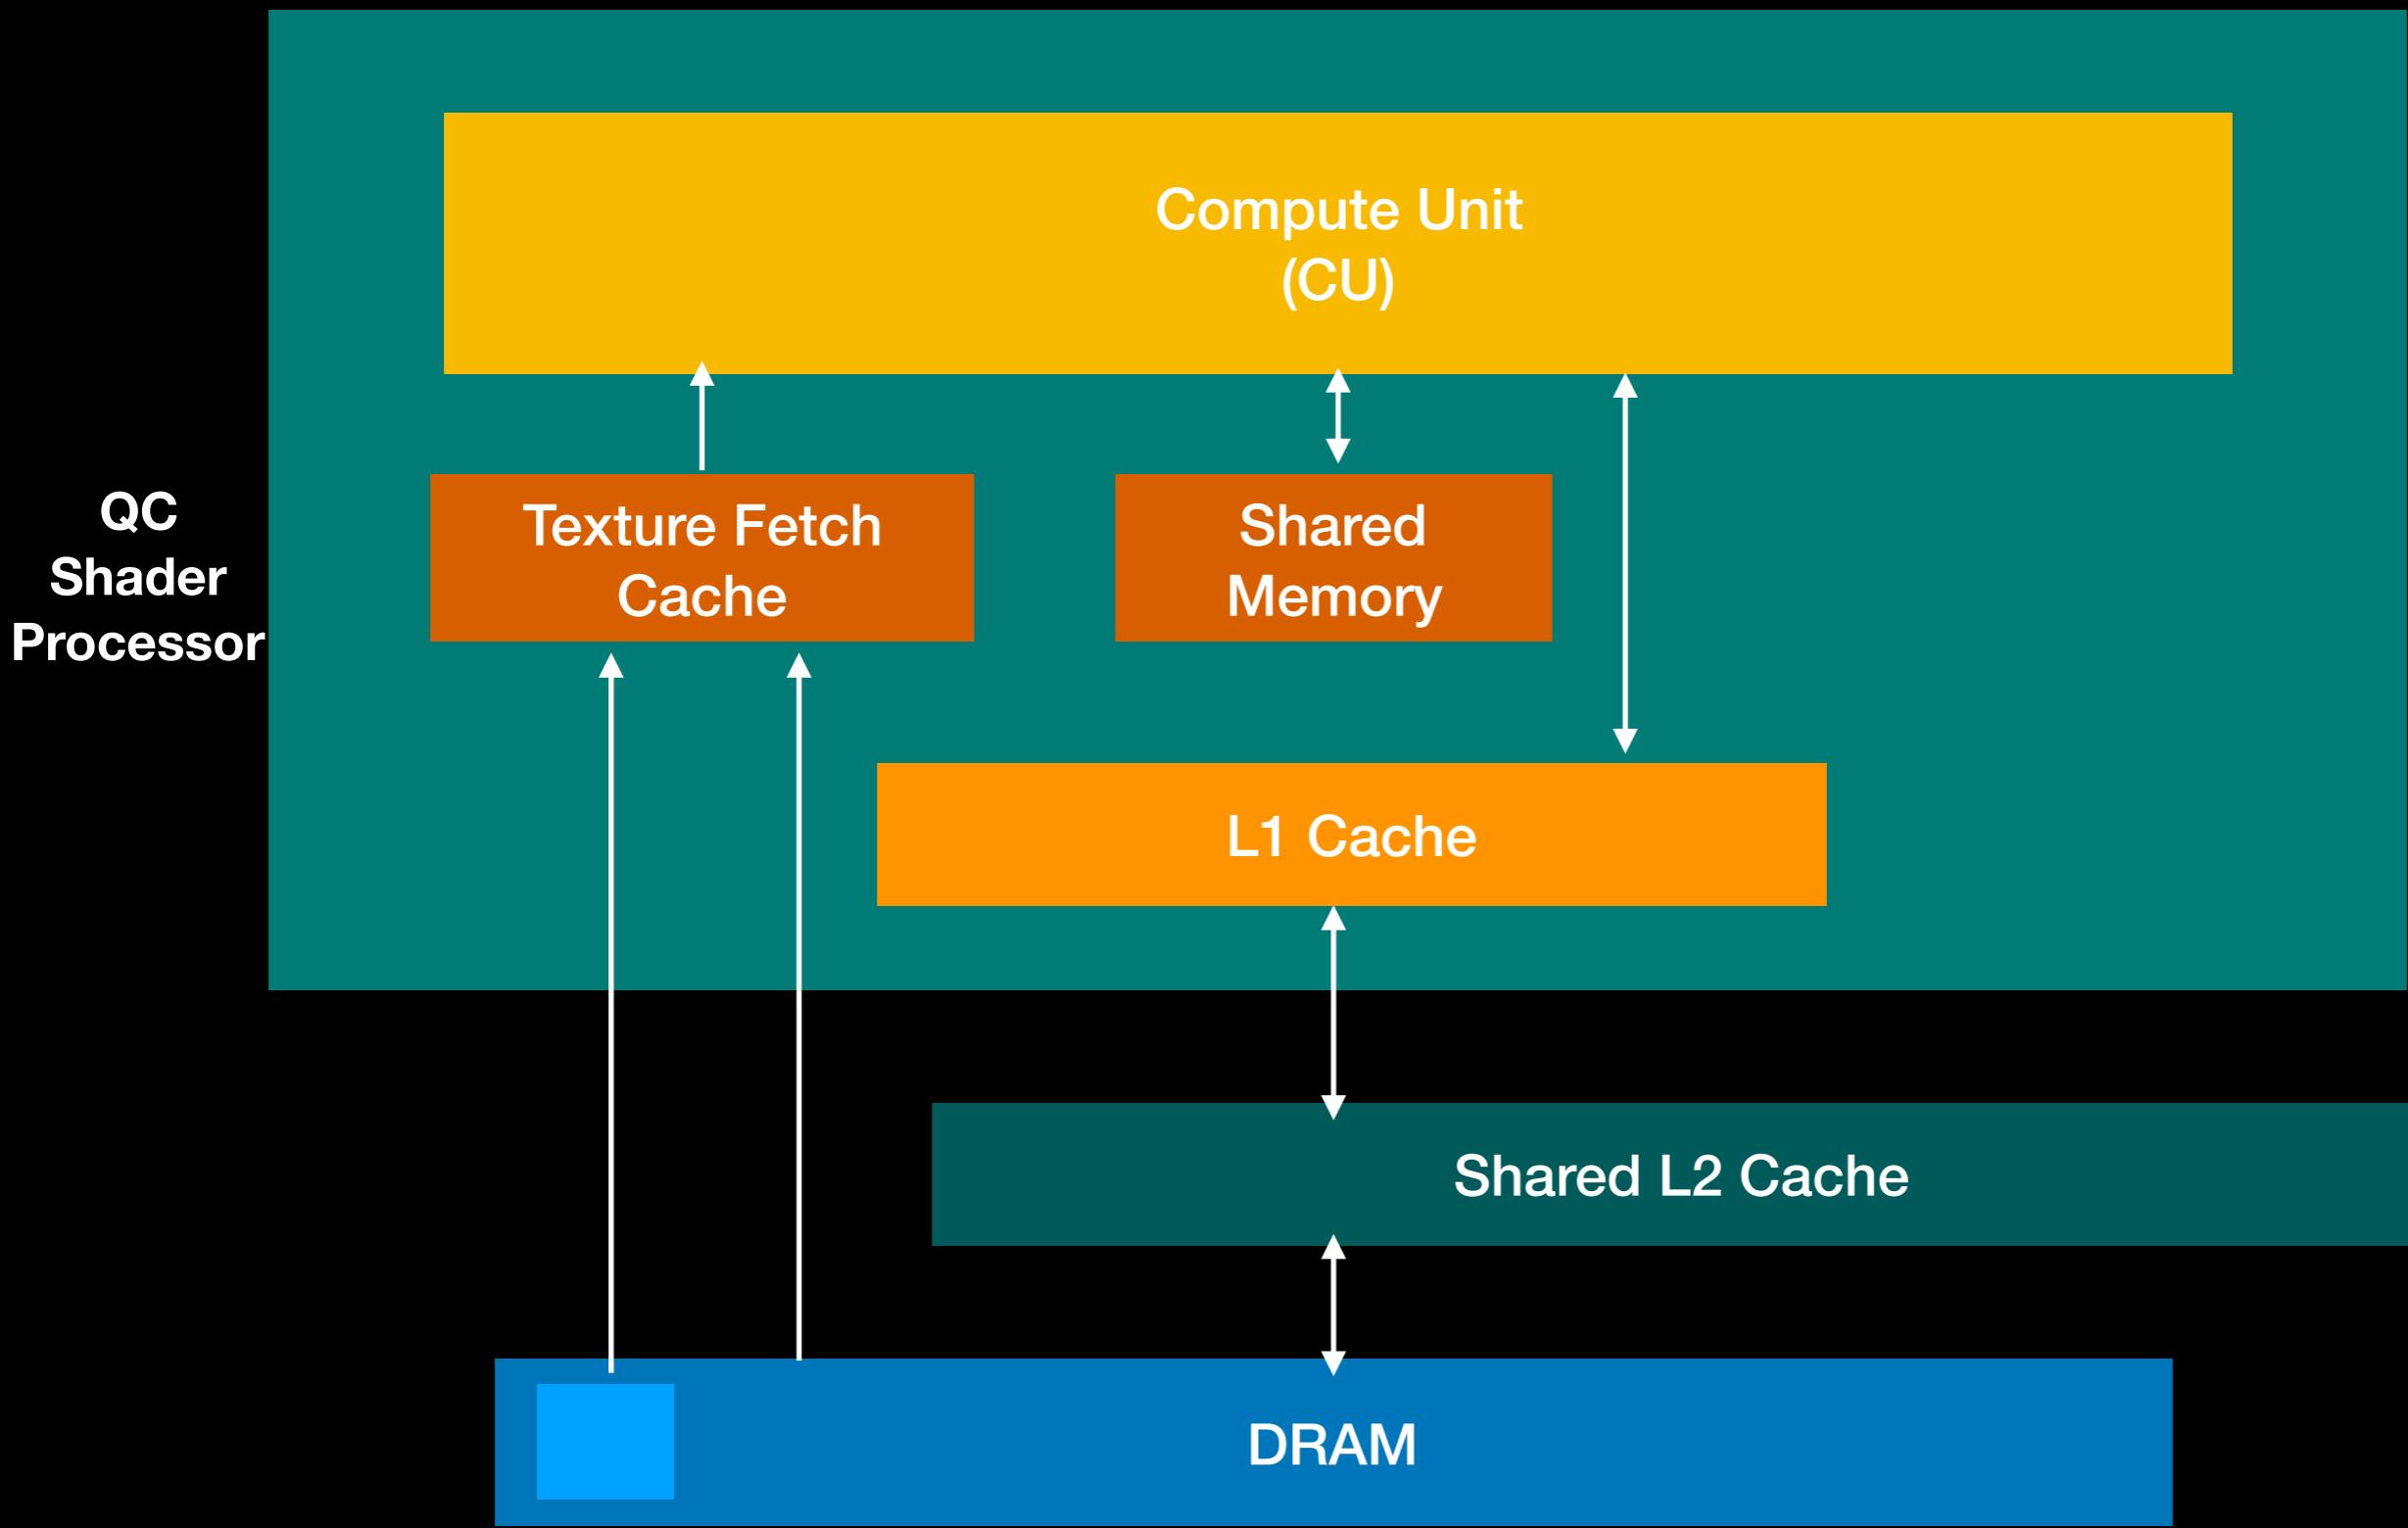
# Shader Processor View



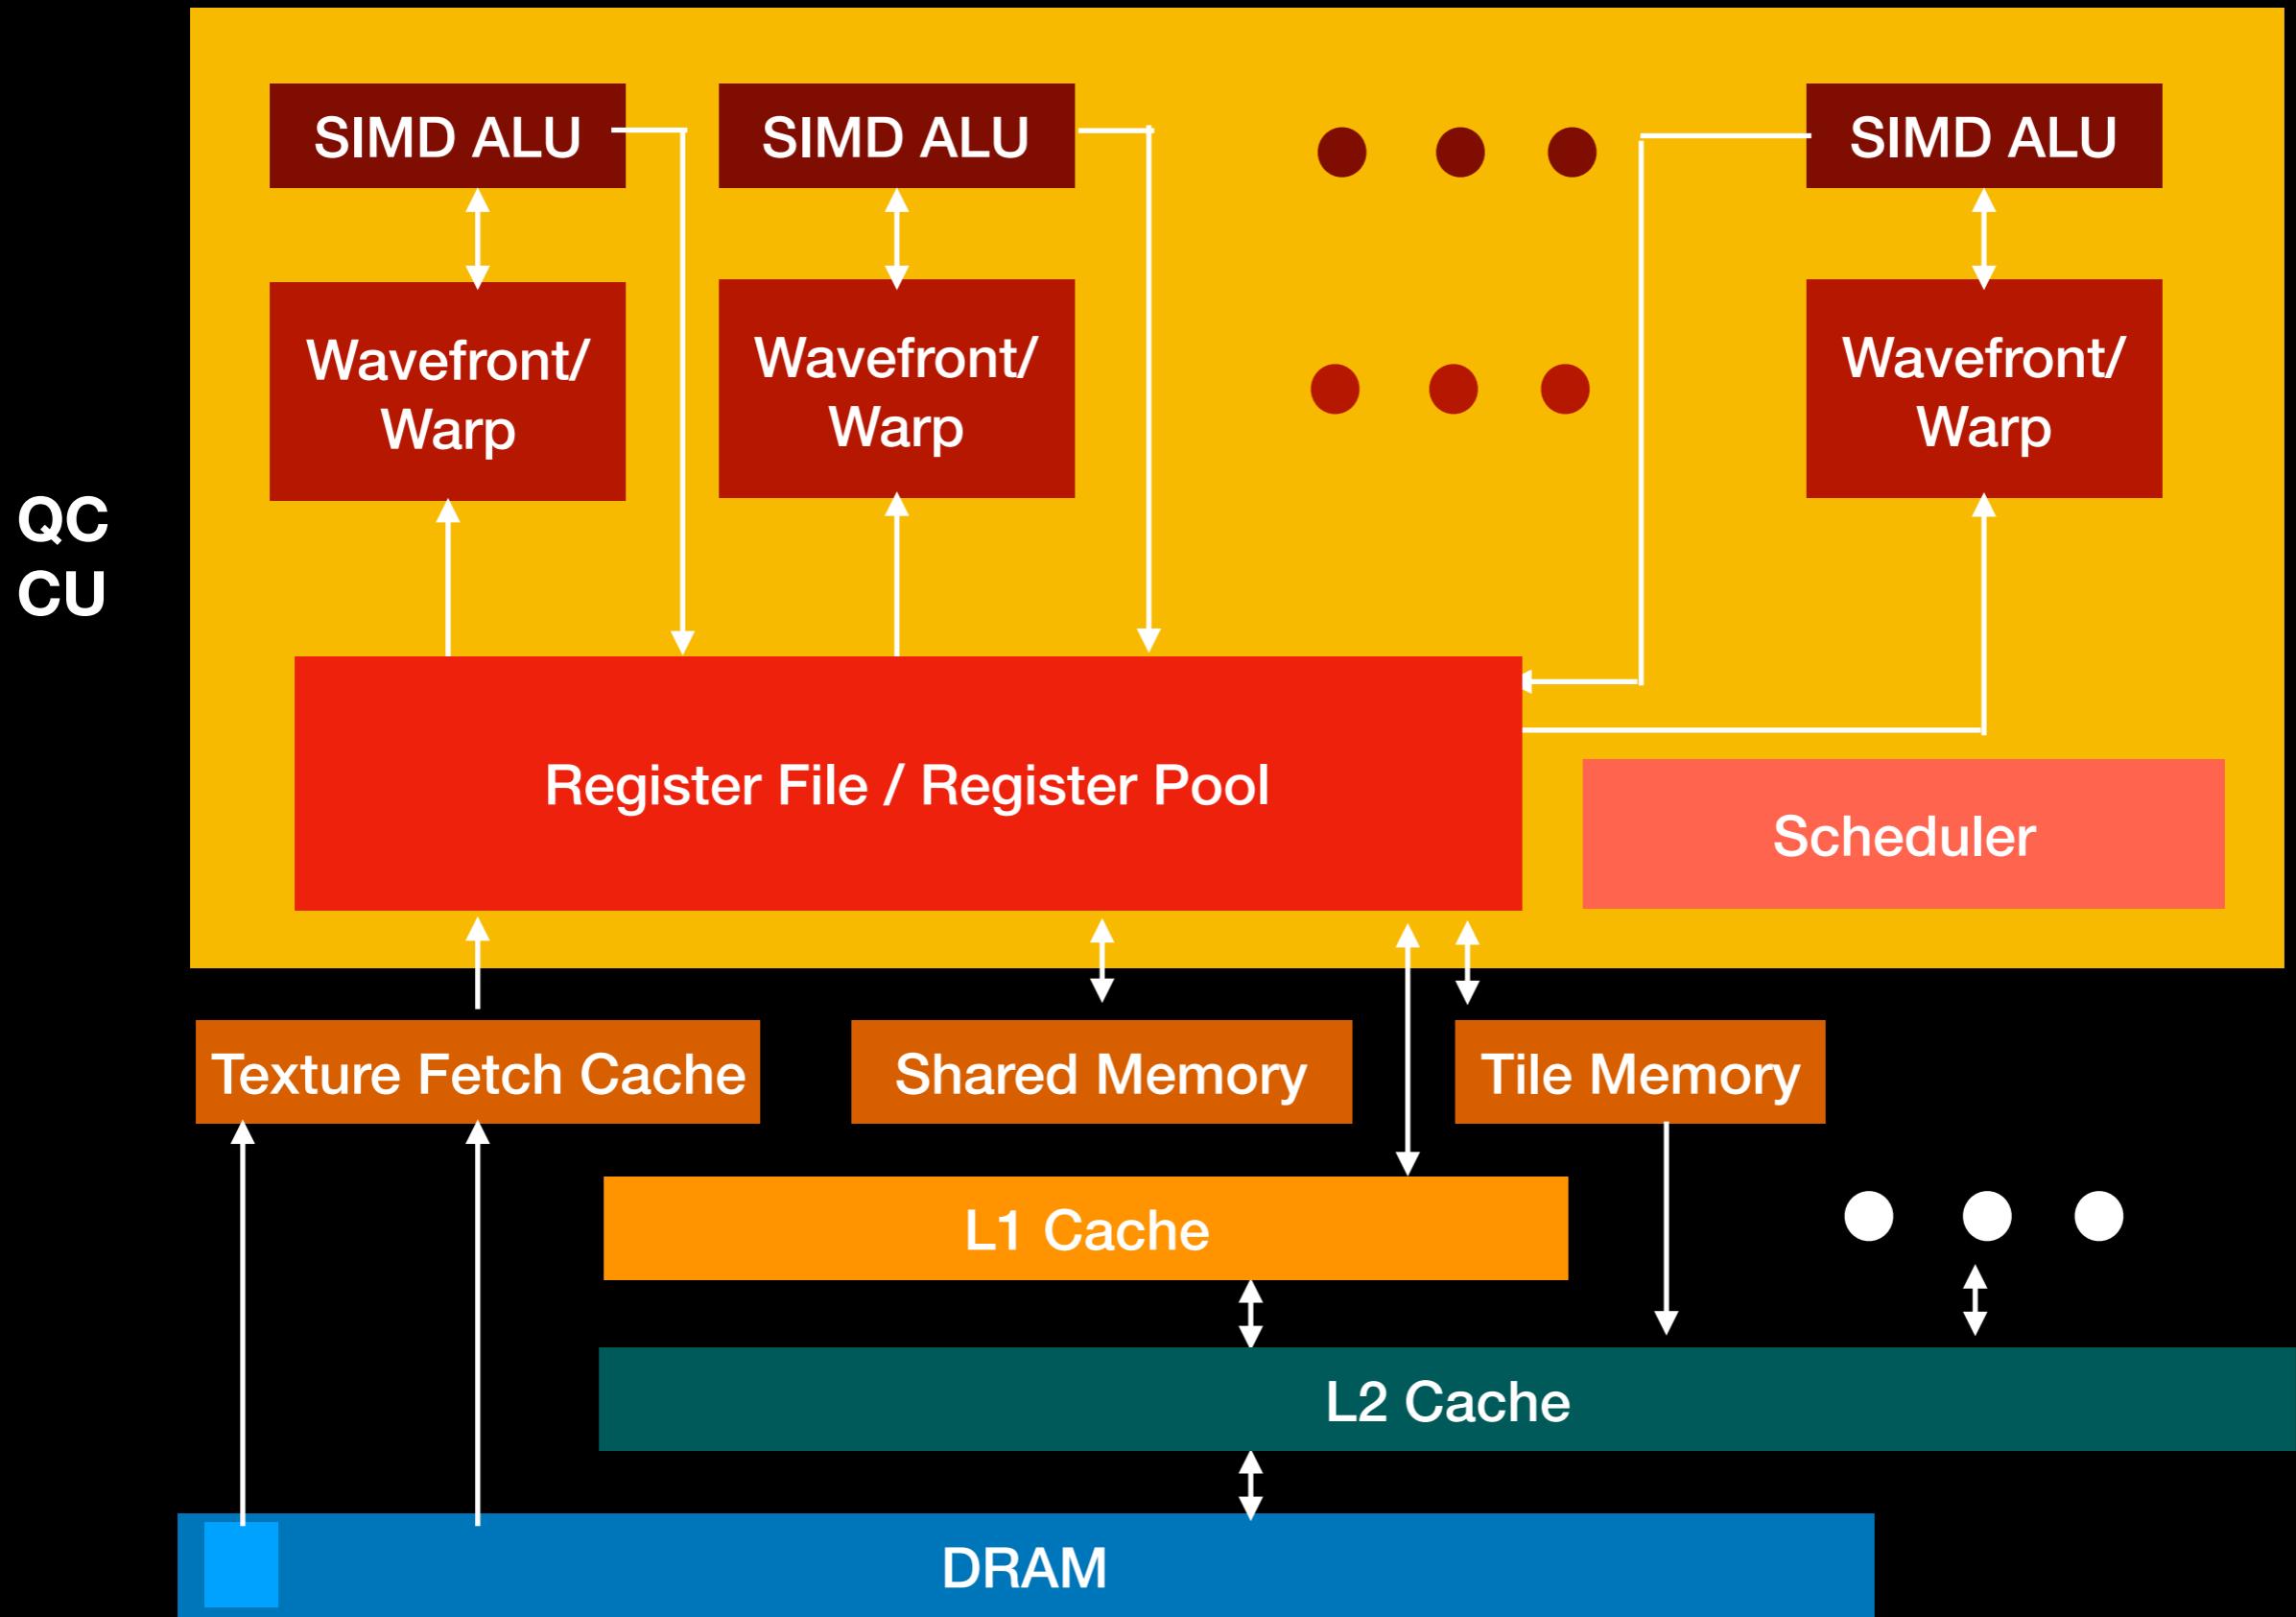
# OpenGL Fragment Shader Path



# OpenGL Compute Shader Path

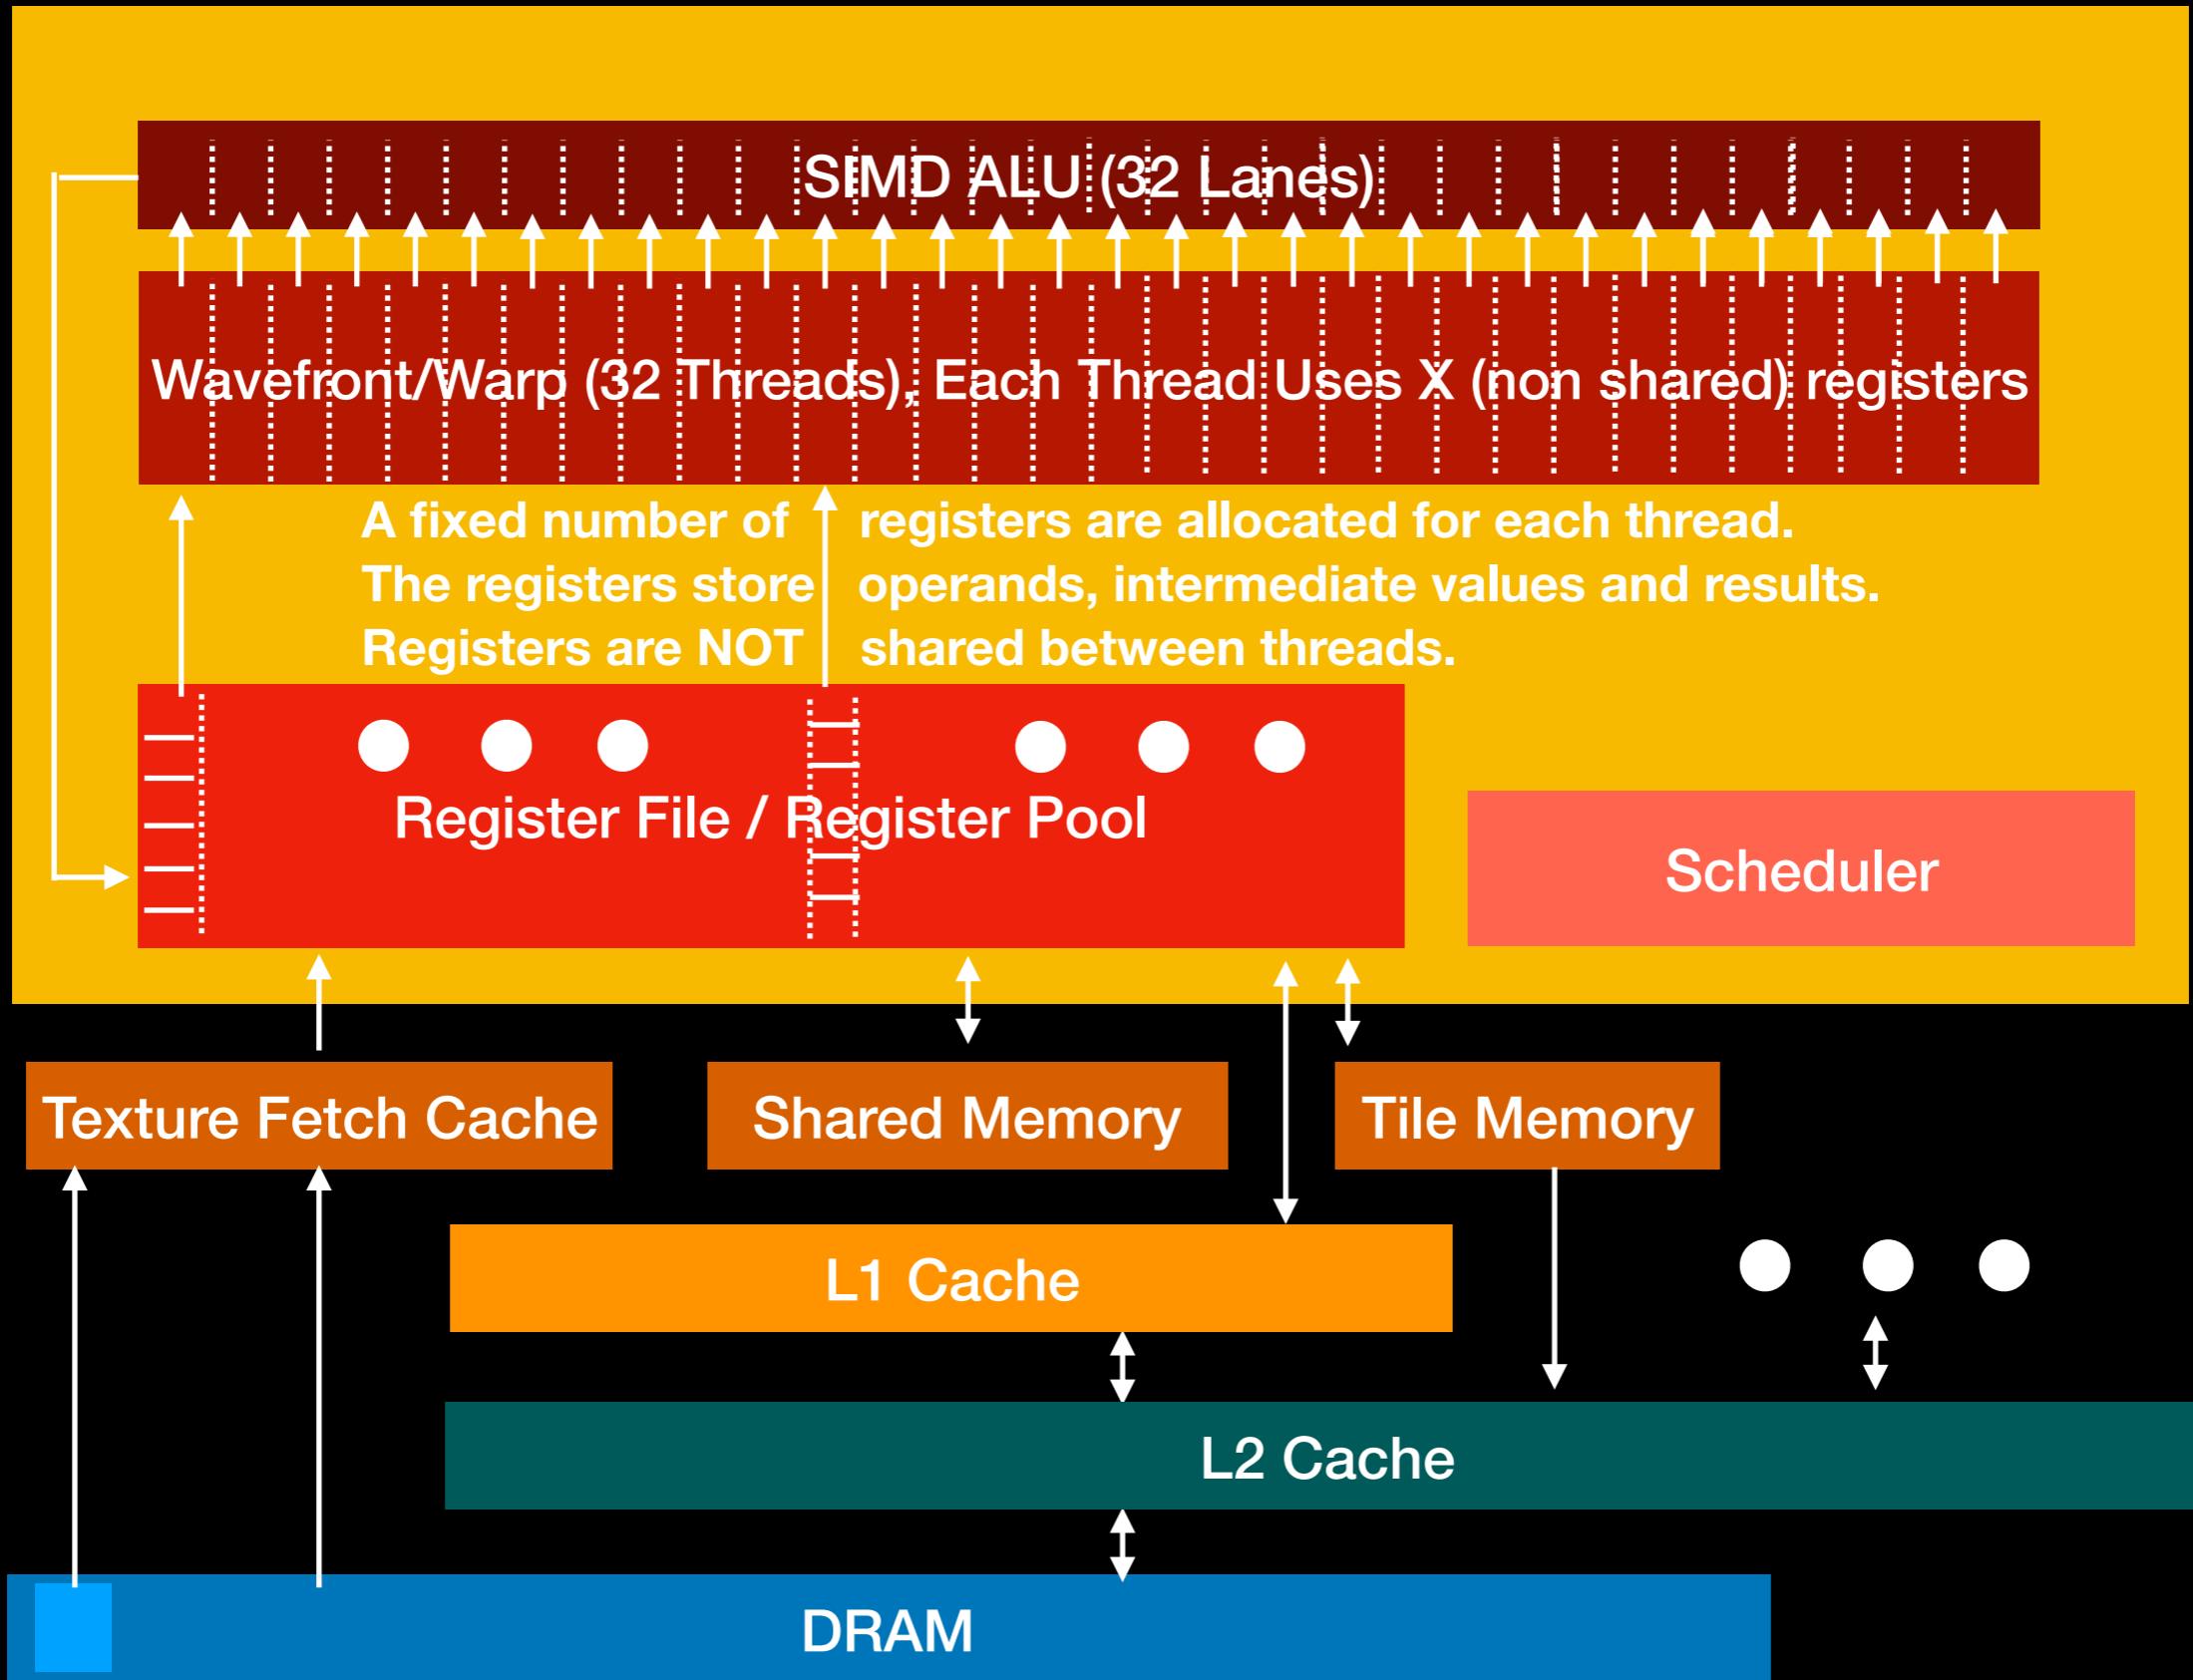


# Compute Unit (CU) View



# ALU and Wavefront Details

ALU  
and  
Wavefront  
Details



# The Details, In Reverse Order

- CU
- Tile and Shared Memory
- Texture Cache

# The CU Details

- The ALU contains 32 (or possibly 64) SIMD lanes.
- Each instruction is executed in SIMD fashion, applying the same instruction to all 32 (or 64) lanes simultaneously.
- During an instruction issue, each lane is assigned one thread from a wavefront.
- Each thread has its own private registers that store operands, intermediate values and results.
- The lanes act as temporary execution slots for threads and are reused across different wavefronts over time.

# The CU Flow

- If a thread in a waveform \*\*“is not ready for processing” - the entire waveform is evicted from the ALU, and a different resident waveform is substituted in it’s place.
- Occupancy describes the availability of substitute resident waveforms.
- Occupancy determines the effectiveness of latency hiding (which will be described shortly).
- \*\*Reasons for not being ready: waiting on a texture fetch, waiting on a memory fetch, waiting on a specialized function unit result (e.g. sin, exp), waiting on a dependency from a prior instruction.

# Tile and Shared Memory

- Tile memory is used by fragment shaders, while shared memory is used by compute shaders.
- \*\*\*Tile memory is hardware driven, and exists “under the hood”.
- Often, programmers / developers are not aware of its existence.
- Shared memory is configurable (via software) by the programmer / developer.
- Both tile memory and shared memory can be read from and written to.

# A Tile Memory Example

- Tile size [GMEM]: 32x32
- Location (depth) [Max Color Attachments]: 8
- Each pixel and location is associated with vec4 [4 Bytes for Int8, 8 Bytes for FP16, 16 Bytes for FP32]
- 32x32x8x4: ~32 KB
- If **total tile memory per SP = 32 KB => 1 tile can exist.**
- If we specify FP16, **the hardware will automatically reduce the tile size**, to say, 16x16.

# Tile Memory

- Tile memory is very close to the CU registers.
- As a result, tile memory is extremely fast.
- Goal: When using a fragment shader, you want to perform all of your compute in an “ALU-register-tile memory” loop.
- Upon completion, tile memory is then written to DRAM.
- The latter minimizes system bandwidth and power usage.

# A Shared Memory Example

- Filter size: 21x21
- Workgroup size (threads): 32x32
- Memory per pixel: 4 Bytes
- => Configure shared memory tile size as:  
 $(32 + 21 - 1) \times (32 + 21 - 1) = 52 \times 52$
- $52 \times 52 \times 4 = \sim 10 \text{ KB}$
- If total shared memory = 64 KB => you can fit 6 workgroups into shared memory.

# Shared Memory

- Shared memory is configured by the user.
- The user specifies the workgroup size and the tile size.
- The workgroup size represents the number of threads associated with the tile.
- Each thread has access to all of the data in the tile.
- For latency hiding reasons, shared memory should be configured so that multiple workgroups result.

# Shared Memory

- In the current example, 6 different workgroups can be realized / running in parallel, using shared memory.
- These workgroups are independent of each other, producing independent wavefronts that are amenable for latency hiding.

# Tile Memory vs Shared Memory

In OpenGL, this translates to:

- When should you use a fragment shader vs a compute shader?

# Fragment Shader vs Compute Shader

- Q: If I wanted to implement XYZ on the GPU, should I use a fragment shader or a compute shader?
- A: What is XYZ more amenable to?

# Fragment Shader vs Compute Shader

- Q: Can XYZ be easily implemented using either shader?
- A: If no, the compute shader will be the shader of choice.
- The compute shader was introduced, to address the implementation needs of various tasks, not amenable using a fragment shader.
- Fragment shaders are targeted for a graphics pipeline.
- Compute shaders turn the GPU from a graphics-centric processor into an explicitly programmable general-purpose parallel processor.

# Shader History - A Brief Segue

- Graphics were originally performed on the CPU.
- However, this proved to be unacceptably slow, for many rendering tasks.
- The GPU (specialized graphics hardware) was introduced, providing parallel compute / vectorization capabilities.
- OpenGL was originally introduced to provide an API to abstract graphics hardware for a fixed graphics pipeline.
- Later, OpenGL added vertex and fragment shaders, to allow a programmable graphics pipeline.

# Shader History - A Brief Segue

- The image / video processing community recognized the compute power of the GPU, and adopted OpenGL's fragment shader for their own needs.
- The fragment shader did an excellent job at various per pixel, or, small neighborhood image / video based tasks.
- However, fragment shaders were restrictive: algorithms requiring large spatial neighborhoods or arbitrary memory access were difficult to implement efficiently.

# Shader History - A Brief Segue

- Other domains also recognized the GPU's compute potential, and faced similar constraints.
- This led to the introduction of the compute shader, and enabled general-purpose GPU computing.
- The compute shader allowed flexible memory access, workgroup cooperation, and arbitrary thread organization.

# Fragment Shader vs Compute Shader

- Assume XYZ be easily implemented using either shader.
- Q: Which shader do we now choose?
- A: When using a fragment shader, does the shader need to frequently (and repeatedly) fetch identical data from DRAM?
- A: Is the overhead for using a compute shader exceedingly high for the task at hand?

# Fragment Shader vs Compute Shader

An additional consideration:

- When writing to DRAM, fragment shaders perform a single “burst” tile write (32x32), after all 1024 results are available.
- In contrast, when writing to DRAM, compute shaders perform 1024 individual writes - when each result becomes available.
- => Compute shaders use more power than fragment shaders.

# Classic Example A

- You are performing an operation using a 21x21 filter kernel - say averaging.
- With a fragment shader, each thread will require its own 21x21 copy of the data.
- Since there is no such thing as shared memory, the SoC hardware quickly gets overwhelmed.

# Classic Example A

- 1) The same data is re-fetched from DRAM numerous times, increasing system bandwidth pressure.
- 2) To address the slowness of the DRAM fetches, the compiler pre-fetches the data by unrolling the averaging / summation loop - increasing register pressure.
- a0) If there are an insufficient number of registers, L1 cache spillover occurs.
- a1) If the spillover is large enough, L1 cache thrashing occurs.
- a2) Because of the increased register pressure, occupancy drops, reducing the effectiveness of latency hiding.

# Classic Example A

- With a compute shader, data is shared amongst the threads in a workgroup.
- 1) Bandwidth pressure is reduced, since each thread does not need an individual copy of the data.
- 2) Loop unrolling does not occur, because pre-fetching is not needed.
- i.e The data is nearby (1-2 cycles vs 100's of cycles).
- 3) Register usage remains low, increasing occupancy and the effectiveness of latency hiding.

# Classic Example A

- The Compute Shader is the clear winner.

# Classic Example B

- You are performing an operation using a 3x3 filter kernel - say averaging.
- With a fragment shader, each thread will require its own 3x3 copy of the data.
- Although there is no such thing as shared memory, the SoC hardware is sufficient for the task.

# Classic Example B

- 1) System bandwidth pressure is greatly reduced, due to the reduced kernel size.
- 2) Loop unrolling is not needed, because the data can be found in the L1 cache lines.
- => Occupancy is high, and latency hiding is effective.

# Classic Example B

- With a compute shader, data is shared amongst the threads in a workgroup.
- The bandwidth pressure advantage is now greatly reduced.
- Occupancy now is similar to the fragment shader, providing no latency hiding advantage.
- **But**, the overhead for shared memory management now becomes significant.
- **But**, the time needed for thread synchronization now becomes significant.

# Classic Example B

- The Fragment Shader is now the clear winner, since it does not have the associated overhead of shared memory.

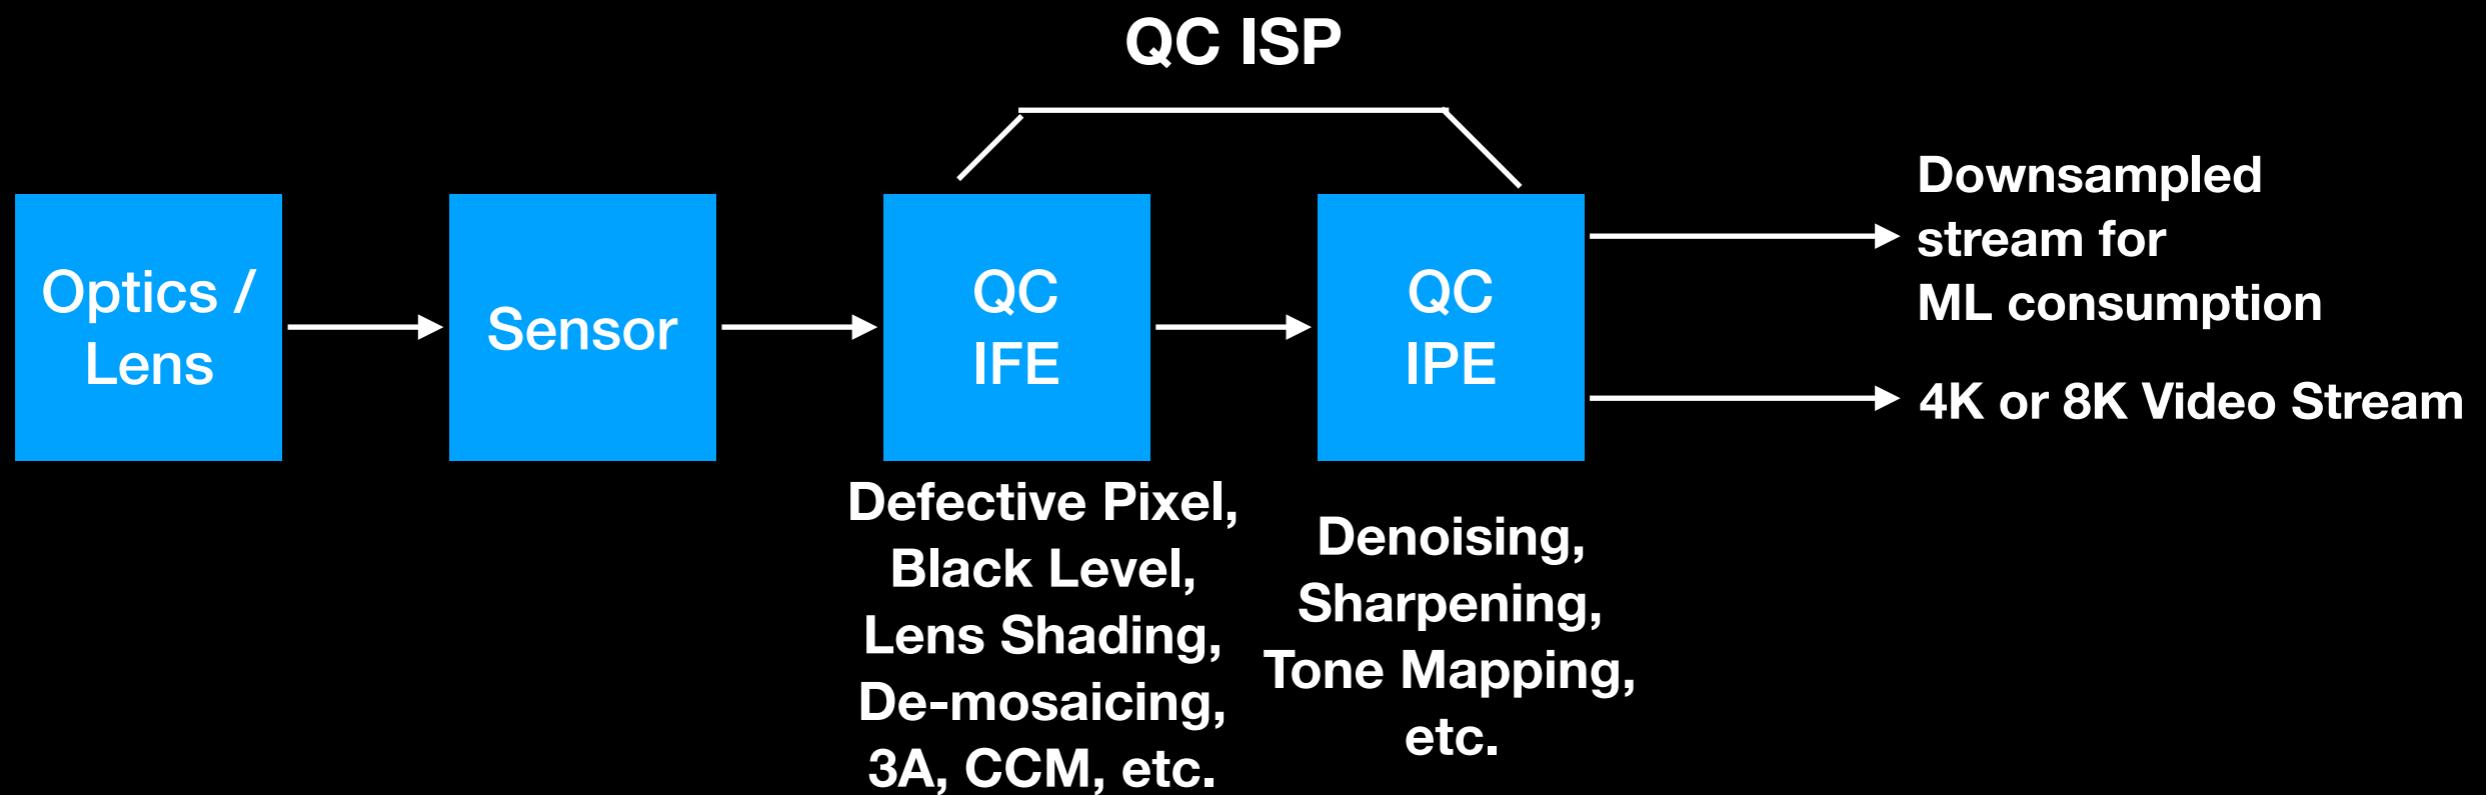
# Classic Example C

- At some point in-between 3x3 and 21x21, a crossover point exists.

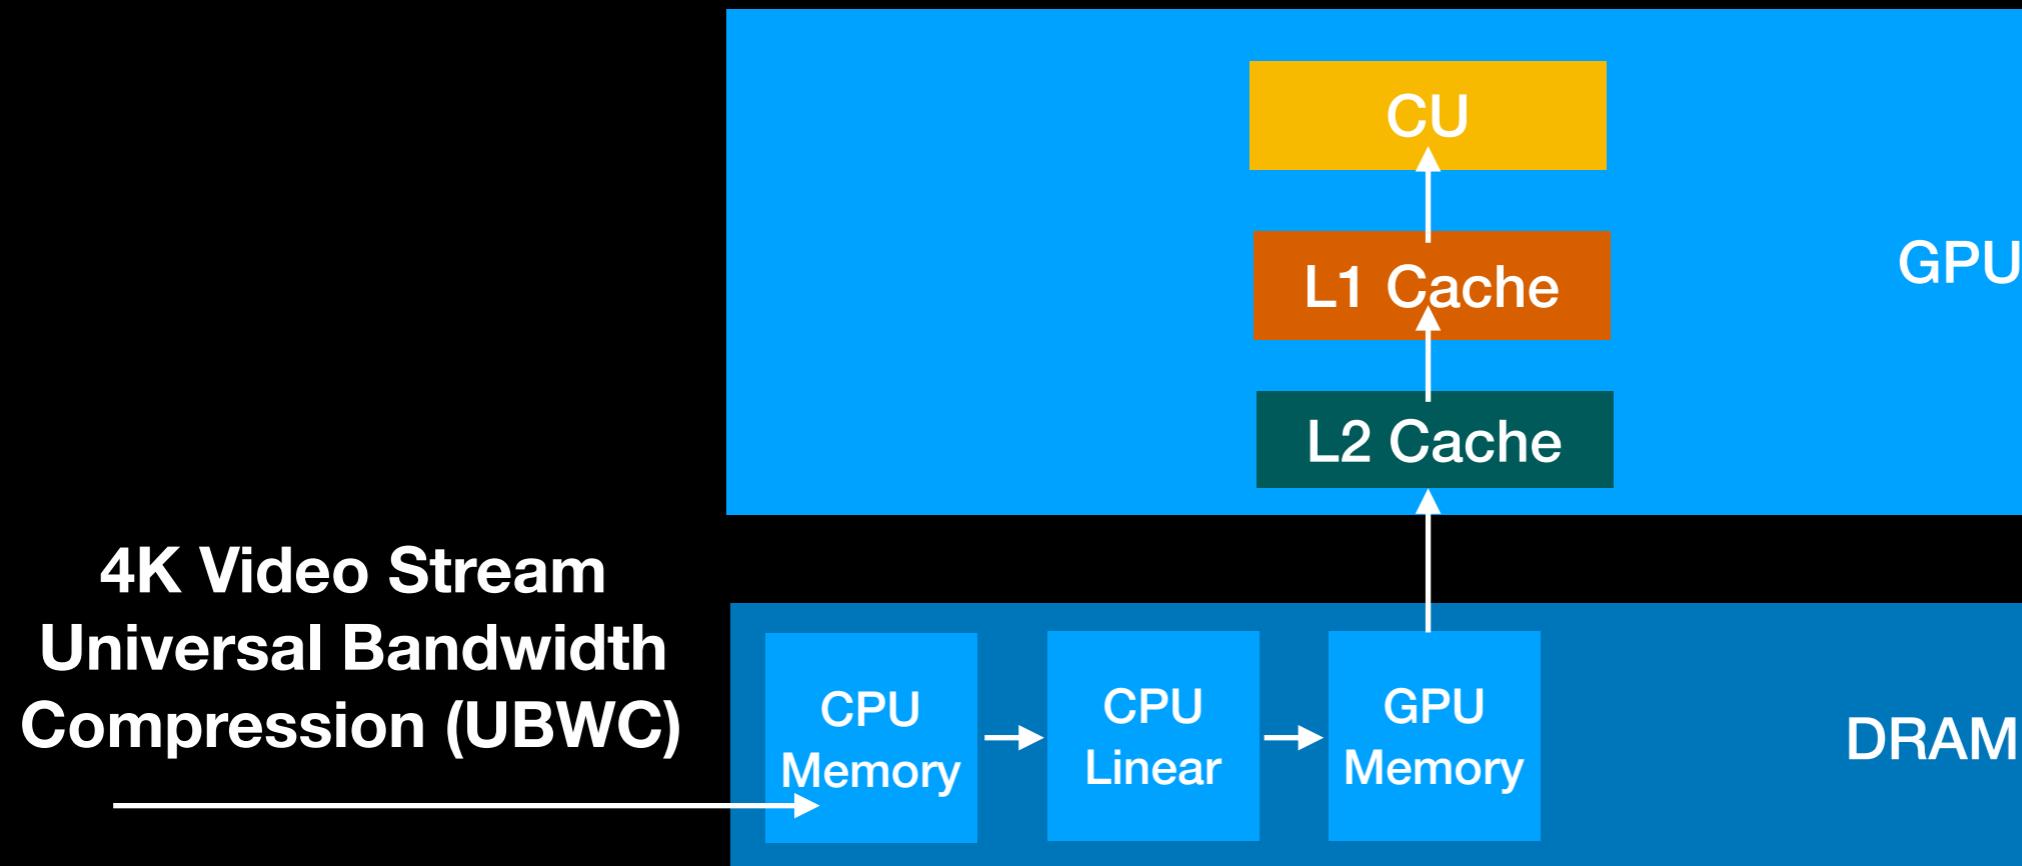
# Texture Cache - Data Fetching

- There are different data fetching paths for the Adreno GPU.
- Traditional (non-native graphics buffer)
- Fast (native graphics buffer)
- We will now illustrate how data migrates from the ISP, using these two paths.

# Signal / Video Stream



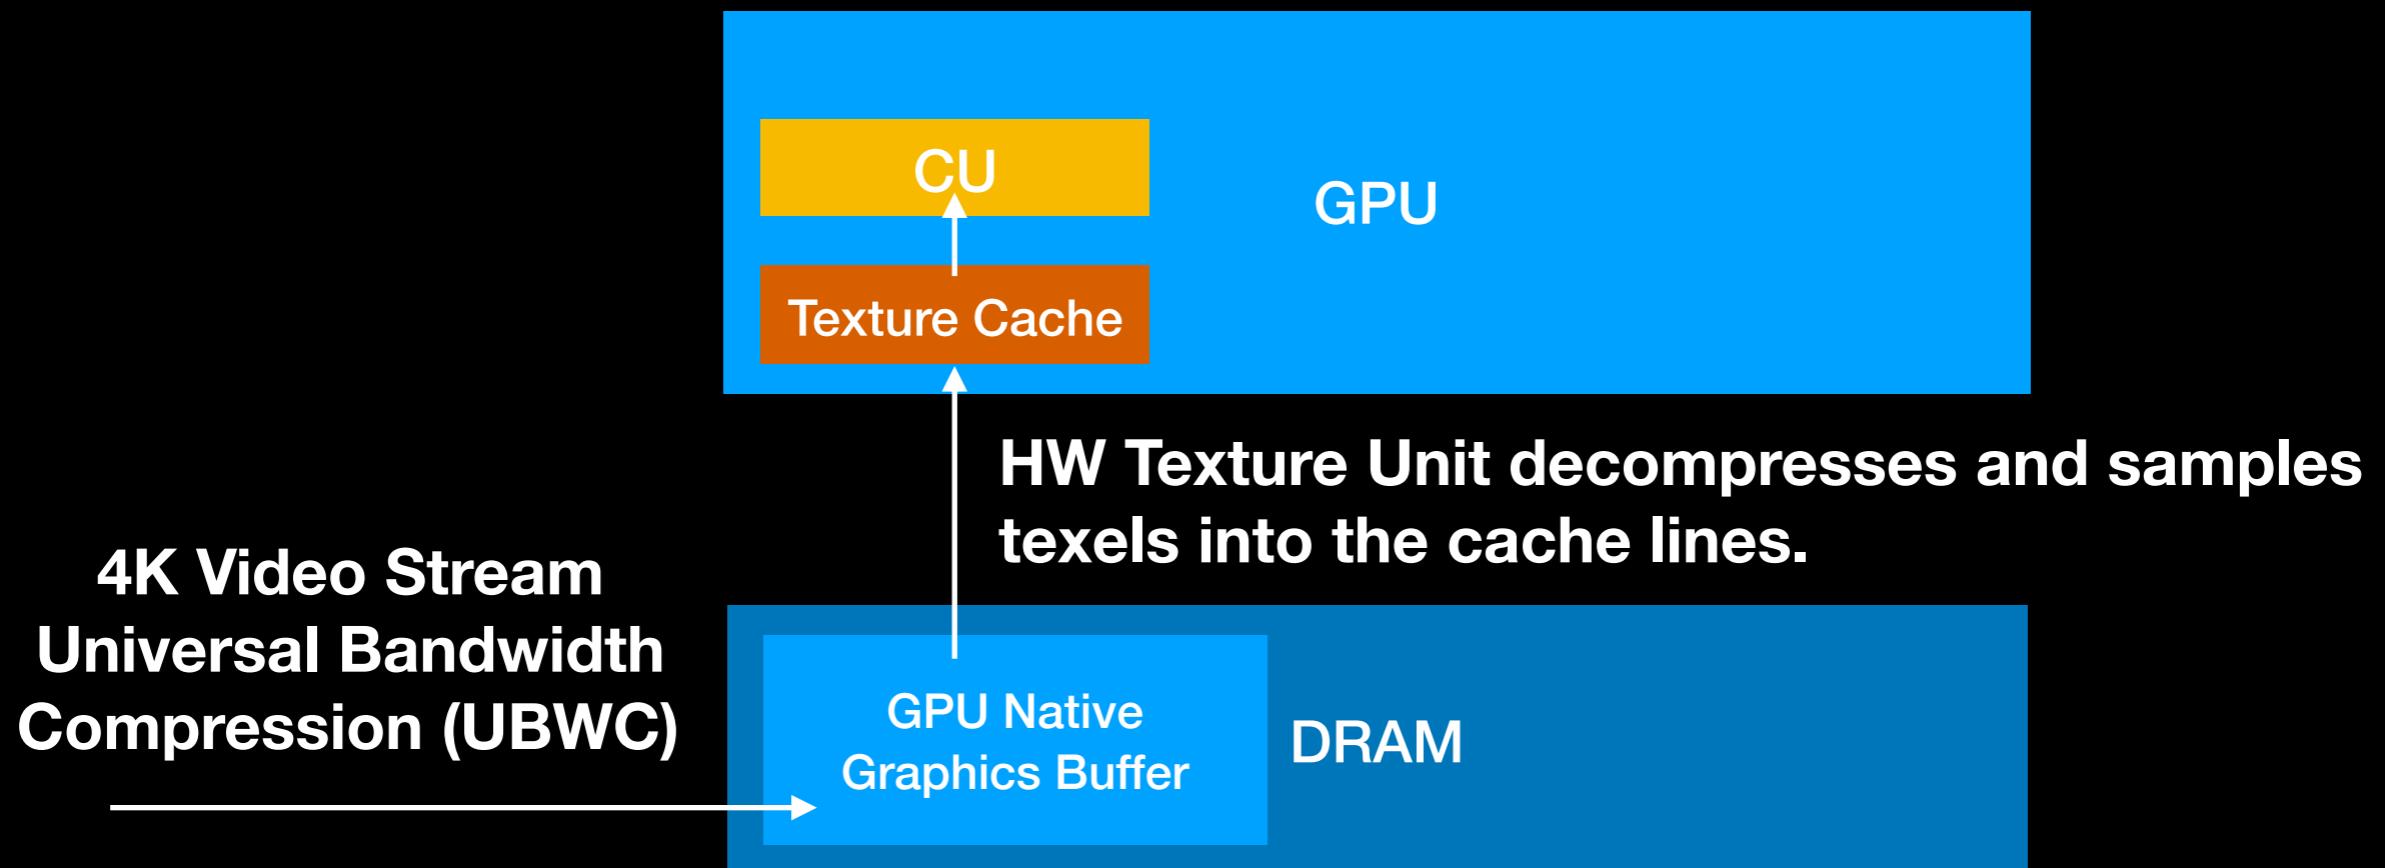
# Video Stream (Traditional): Non Surface Path



Multiple copies occur.

- 1) The **UBWC tiles** are written to **cpu memory**.
- 2) The tiles are decompressed, and written to **standard linear format**.
- 3) The linear format data is written to **GPU specific memory**.
- 4) The data migrates to the CU through the L2 and L1 caches.

# Video Stream (Fast): Surface Path



- 1) A surface / surface texture is defined, to create the native graphics buffer.
- 2) A hardware texture unit (TU) samples the data directly into the texture cache by filling the 64 Byte cache lines.

No copy / memcpy is involved.

# Profiling

- We now have a solid understanding of the Adreno GPU.
- This understanding provides a solid foundation to address the next issue: power and performance.
- Specifically, what are the key idea(s) for optimizing / maximizing this criteria?

# Profiling

- The ALU is the fundamental compute engine in the GPU.
- The objective is to fully utilize all of the ALU's.
- However, this may not be occurring because the ALU's are stalled, waiting for data.

# Profiling

Where does the data need to reside?

- Registers

Where could the data reside?

- Texture Cache
- Tile Memory, Shared Memory
- L1 Cache
- L2 Cache
- DRAM

# Profiling

- How long does it take to move the data into the registers?  
(Ballpark / back of envelope approximations)
- From Texture Cache: (let's assume) 1 clock cycle
- From Tile or Shared Memory: ~1-2 clock cycles
- From L1 Cache: ~2-3 clock cycles
- From L2 Cache: ~10's clock cycles
- From DRAM: ~100's clock cycles

# Profiling

- Hence, it pays to have the needed data, as close to the registers as possible.
- The search for the data begins locally and moves outward
  - texture fetch cache, shared memory, L1, L2, and finally, DRAM.
- Profiling parameters like ALU utilization, texture fetch stalls, L1 hit rate, etc. provide valuable insight into potential bottlenecks.

# Profiling

- How large are the different memory units?
- Memory units are kept small on mobile devices, to reduce power consumption.
- The scaling might look something like:

# Profiling

- Register: 16 B
- Register File (Main reason occupancy is limited): 64 KB
- Cache Line: 64 B
- Texture Cache (Read Only): 8 KB
- Tile memory: 32-64 KB, Shared Memory: 64 KB
- L1 Cache: 16-64 KB
- L2 Cache: 512 KB-8 MB
- DRAM: 8 GB

# Profiling

- Understanding memory usage is important.
- Memory becomes a problem, when there is “not enough” where it is needed.
- Classical Example A (fragment shader) provides one such example.

# Two Design Alternatives

- For a fixed and known workflow, one could scope the problem, and construct the optimal configuration to ensure that the ALU's have a continuous stream of usable data.
- If one needed a robust workflow that supported a variety of different use case scenarios, the latency hiding trick can serve as an alternative.

# Latency Hiding Trick

- The ALU SIMD is fed by a **wavefront / wave** consisting of 32 (or 64) threads with associated registers and data.
- Suppose there was a thread in the wavefront that did not have the needed data, in one (or several) of it's registers.
- If this was the only wavefront, the ALU would need to wait, until the data arrived. **BAD!**
- Suppose there was another wavefront that was “ready to go”.
- The “ready to go” wavefront could (then) be swapped in, resulting in little ALU downtime = **LATENCY HIDING!**

# Latency Hiding Trick

- Ideally, if there are  $M$  ALU's, we would like to have  $N$  wavefronts, where  $N$  is significantly greater than  $M$ .
- That way, we could greatly exploit latency hiding.
- However, the number available wavefronts depends on the chip resources - such as the number of total registers.

# Latency Hiding Trick

- i.e. The register file / register pool is finite.
- If each thread uses a large number of registers, fewer wavefronts would be possible, negating the latency hiding trick.
- Also, when using a compute shader, the shared memory should be configured to encourage multiple workgroups.
- This is because the threads in each workgroup will (then) be independent of the threads in a different workgroup.

# In The Limit

- Suppose that a single wavefront consisting of 32 (or 64) threads, used up the entire register pool.
- Then, latency hiding would be impossible, because there would be no other wavefront to “swap in”.
- In profiling jargon, occupancy is low / near zero.
- Ideally, you want to have a lot of available wavefronts that are ready to go / ready to “swap in” = high occupancy.

# In The Limit

- Suppose that the shared memory in your compute shader was configured so that a single tile consumed all of the shared memory.
- Although you now have a large workgroup with many threads resulting in many resident wavefronts, the resident wavefronts are not independent.
- As a result, latency hiding is ineffective, because a problem existing in one wavefront will be replicated in every other wavefront.

# Free Gains To Reduce Hardware “Pressure”

- Tried and true tricks can be applied to ameliorate issues that prevent latency hiding - assuming the output quality remains acceptable.

The underlying theme: process less data

- 1) Downsample, process the reduced data, upsample
- 2) Decrease the precision, where possible: FP32 -> FP16 -> INT8

Both allow the creation of additional independent wavefronts, increasing the effectiveness of latency hiding.

# Summary

- The goal is to fully utilize the ALU's on a GPU.
- However, this can only be achieved, if the ALU's are fed a continuous stream of data.

At this point, there are 2 paths.

- 1) Optimize the pipeline for the given use case, to make this statement as true as possible.
- 2) Employ the latency hiding trick.

# Summary

- Fragment shaders and compute shaders both have utility for different cases.
- Understanding the strengths and weakness of each, arms you with the necessary knowledge to make the right choice.